

# Lecture #16 - Adaptive Quadrature and Load Balancing

AMath 483/583

# Announcements

- **Homework #4** — Due Next Friday @ 5:00pm (Last Day of Class)
  - short, but conceptually challenging
  - (i.e. not something you can do at the last minute)
  - clarification in this lecture
- Today:
  - MPI on multiple machines
  - Return to OpenMP — “nested parallelism”

# Quick Aside

- MPI is for distributed memory environments

```
$ mpiexec -n 4 -H [server list] ./myprog
```

- Can also create a hostfile:

```
# my_hostfile  
foo.example.com  
bar.example.com slots=2  
foobar.example.com slots=4 max-slots=4
```

# Quick Aside

- MPI is for distributed memory environments

```
$ mpiexec -n 4 -H [server list] ./myprog
```

- Can also create a hostfile:

```
# my_hostfile  
foo.example.com  
bar.example.com slots=2  
foobar.example.com slots=4 max-slots=4
```

Two-core machine

Quad-core machine - ensure no more  
than 4 processes run here

# Disclaimer

- Hardest part is setting up computers / ssh-keys
- In following demo
  - americano.amath.washington.edu can send requests to mocha.amath.washington.edu
  - not vice versa (?)


# Demo

hello-distributed.c


# On Homework #4

- **Clarify some misconceptions:** use the *strategy* from `shift.c`, not the same algorithm

```
double* u;
```



```
double* up1;
```




# On Homework #4

```
up1[i] = up[i] +  
        nu*(up[i-1] - 2*up[i] + up[i+1])
```

```
double* u;
```



```
double* up1;
```





# On Homework #4

Solving the periodic domain problem.

$$\text{up1}[N-1] = \text{up}[N-1] + \text{nu} * (\text{up}[N-2] - 2 * \text{up}[N-1] + \text{up}[0])$$

double\* u;

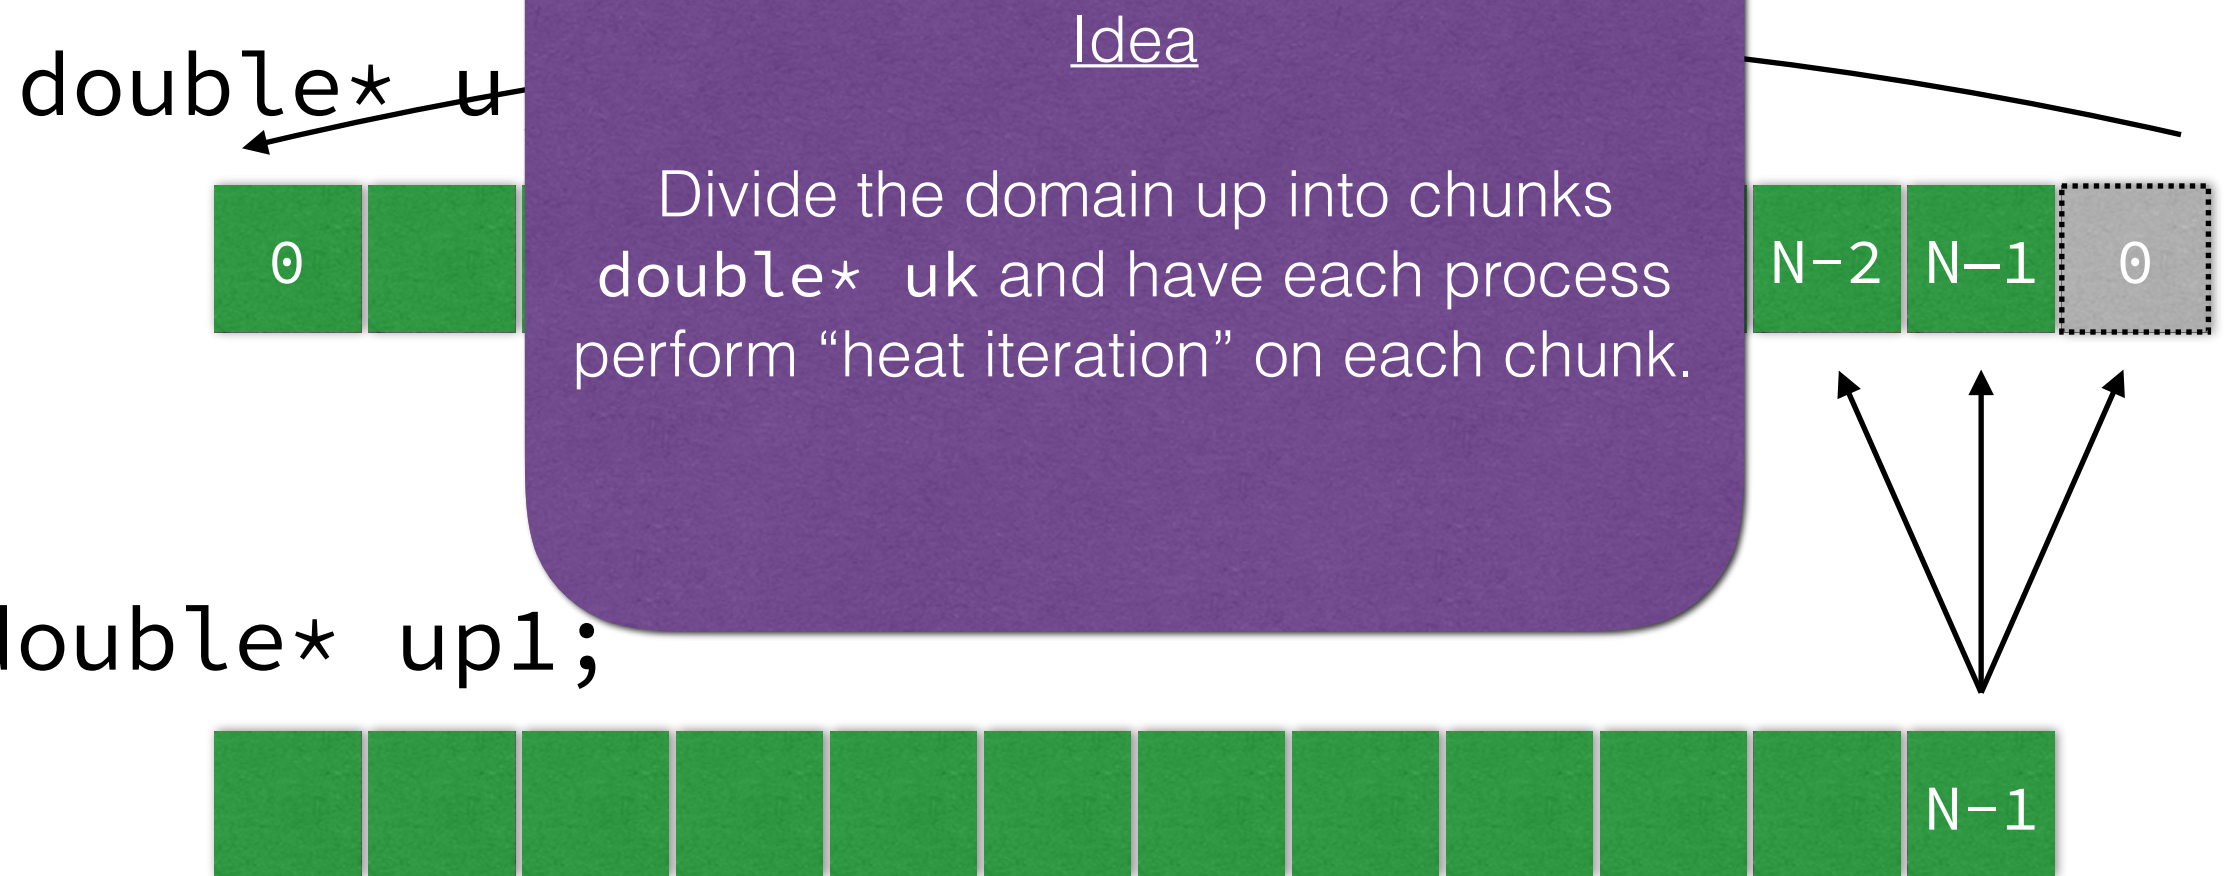


double\* up1;

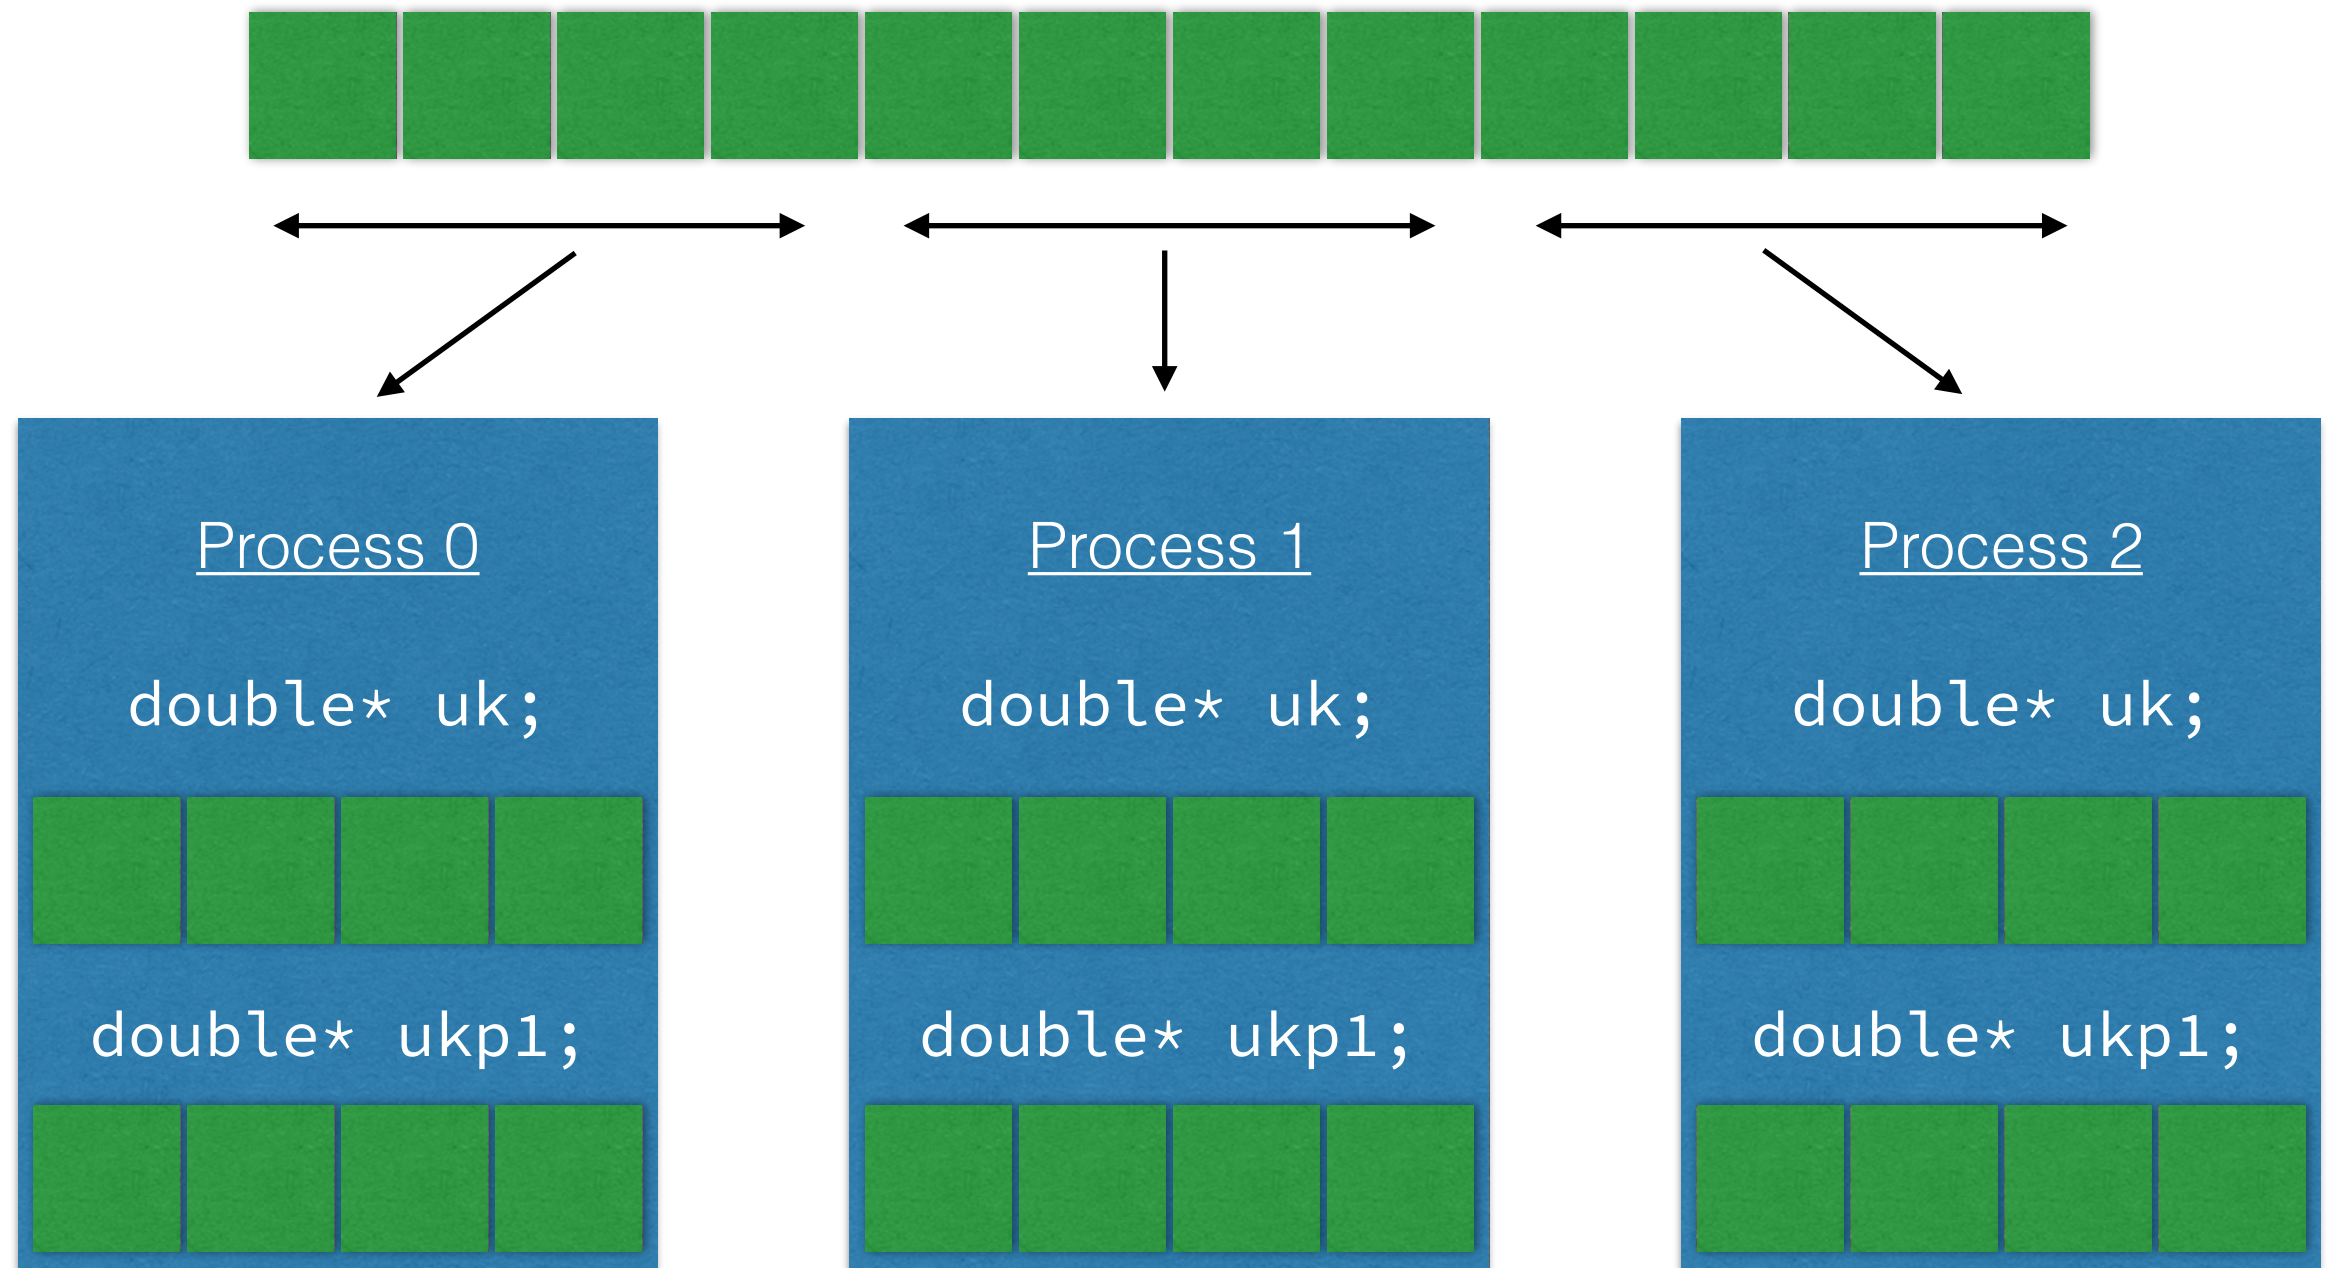


# On Homework #4

$up1[N-1] = up[N-1] + up[0]$


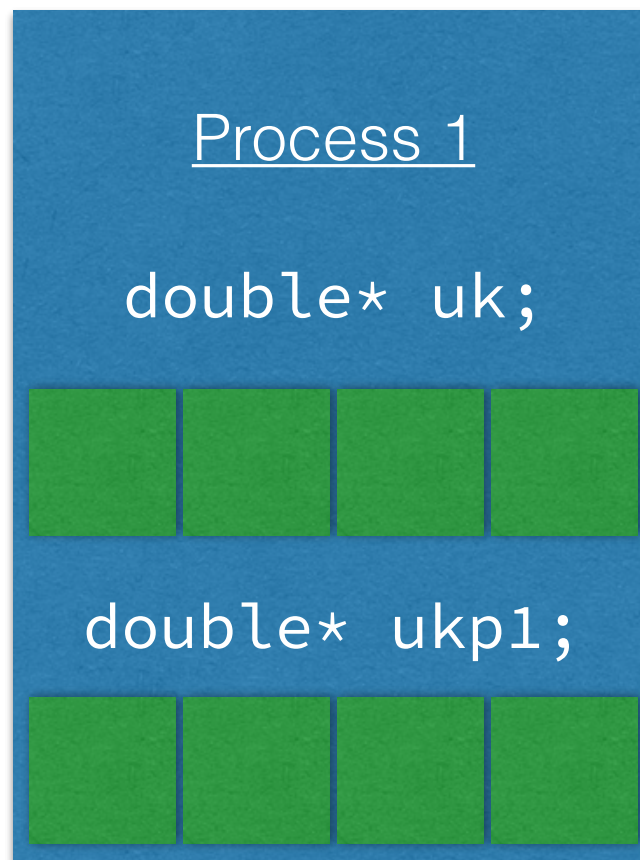
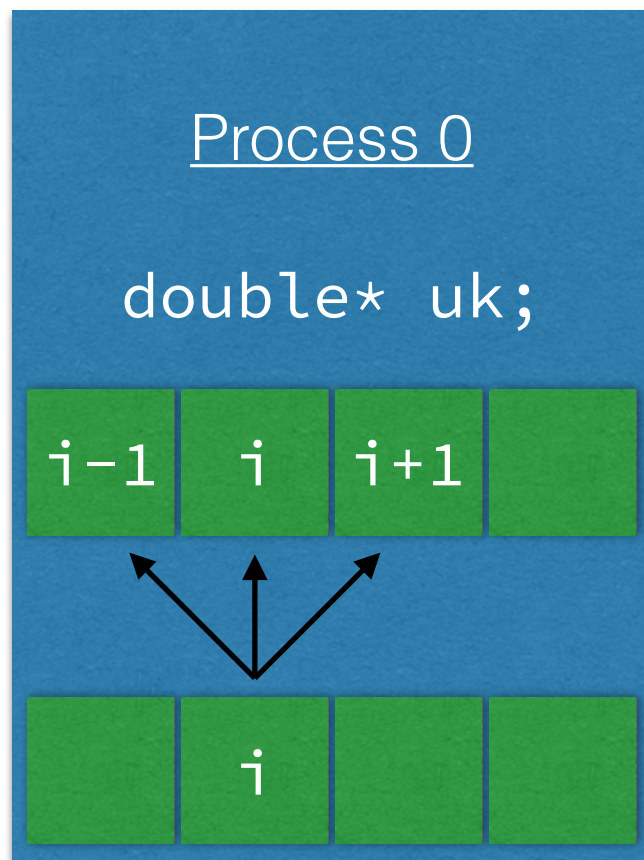


# On Homework #4



# On Homework #4

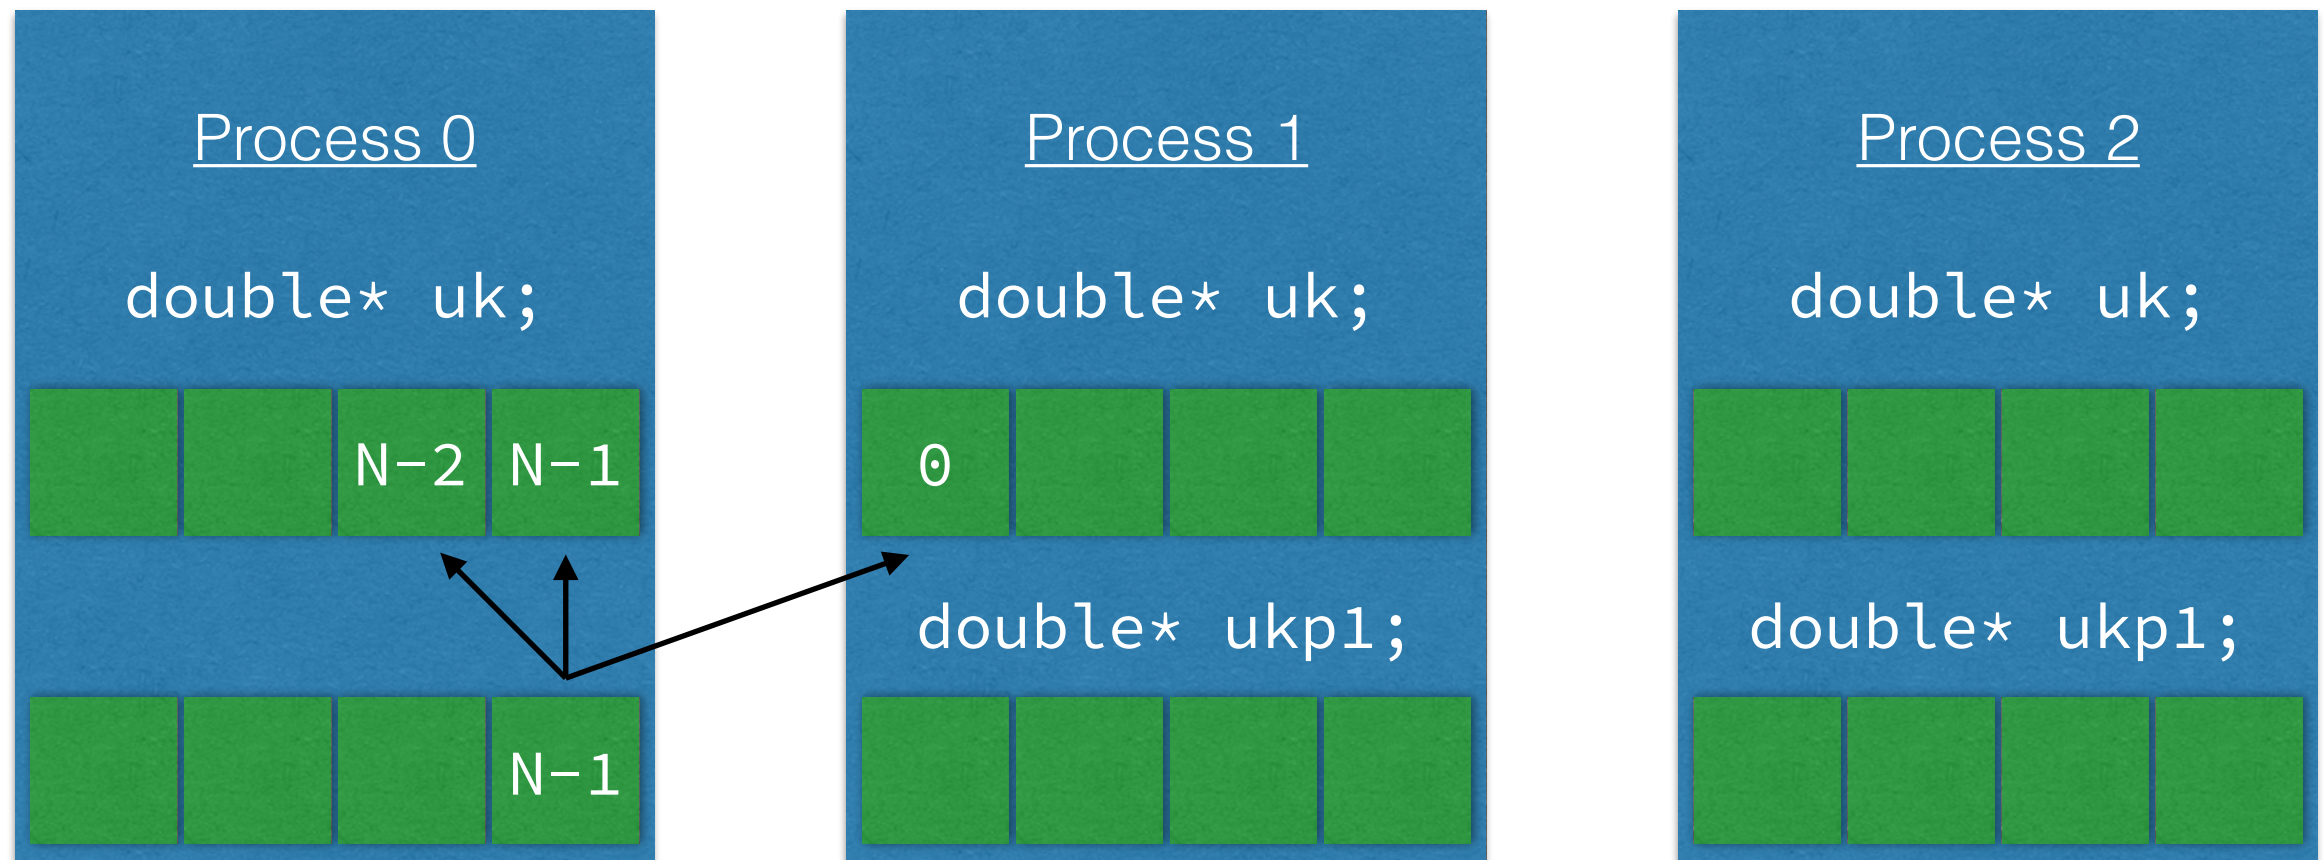
$$\text{ukp1}[i] = \text{ukp}[i] + \text{nu} * (\text{ukp}[i-1] - 2 * \text{ukp}[i] + \text{ukp}[i+1])$$





# On Homework #4

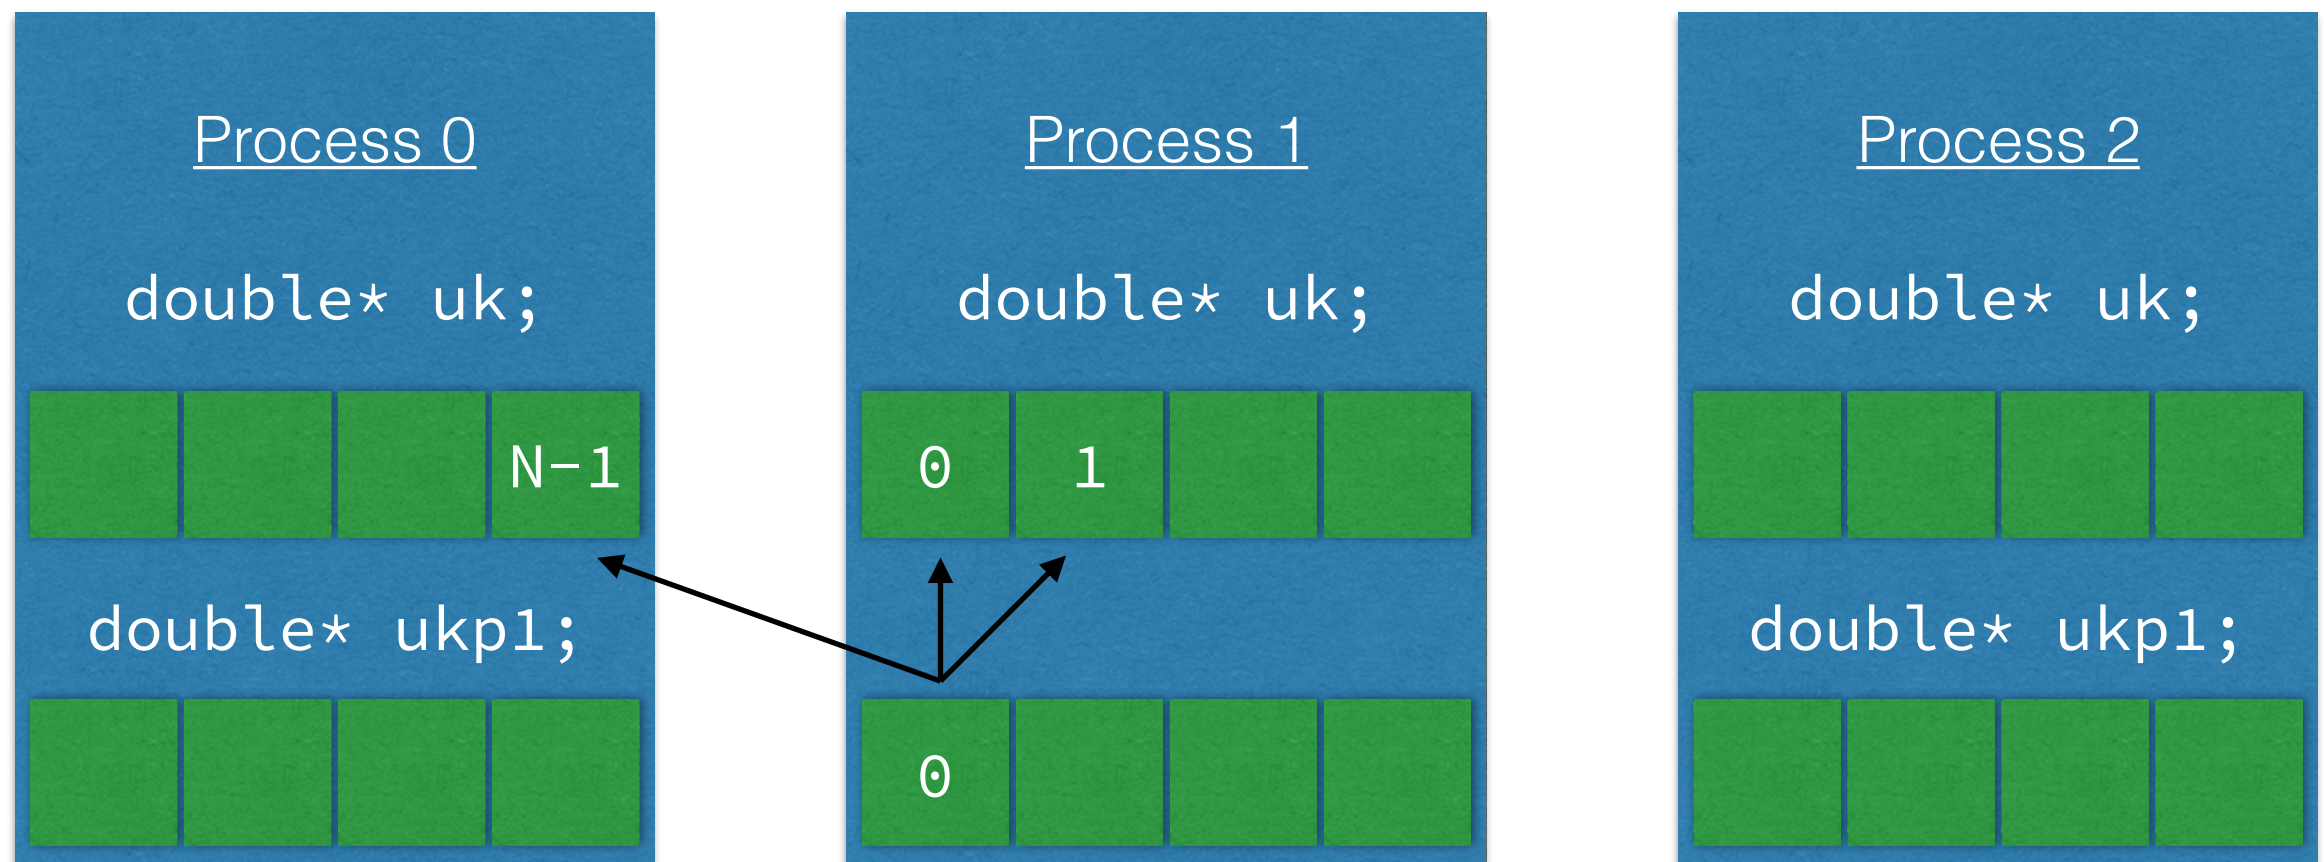
- **Uh oh.** Need to get uk information from neighboring processes. (Send / Recv)



# On Homework #4

- **Same Issue.** Left side of process need info from left proc. Right side of process needs info from right proc.

*(Hint: multiple ghost cells. More than one. More than two?)*



# Adaptive Quadrature


Recursion and Nested Threading

# Adaptive Quadrature

- **Problem:** approximate

$$\int_{-2}^4 e^{-\beta^2 x^2} + \sin(x) dx = \left[ \frac{\pi}{2\beta} \operatorname{erf}(\beta x) - \cos(x) \right]_{-2}^4$$


- where erf is the error function.
- beta = 10:





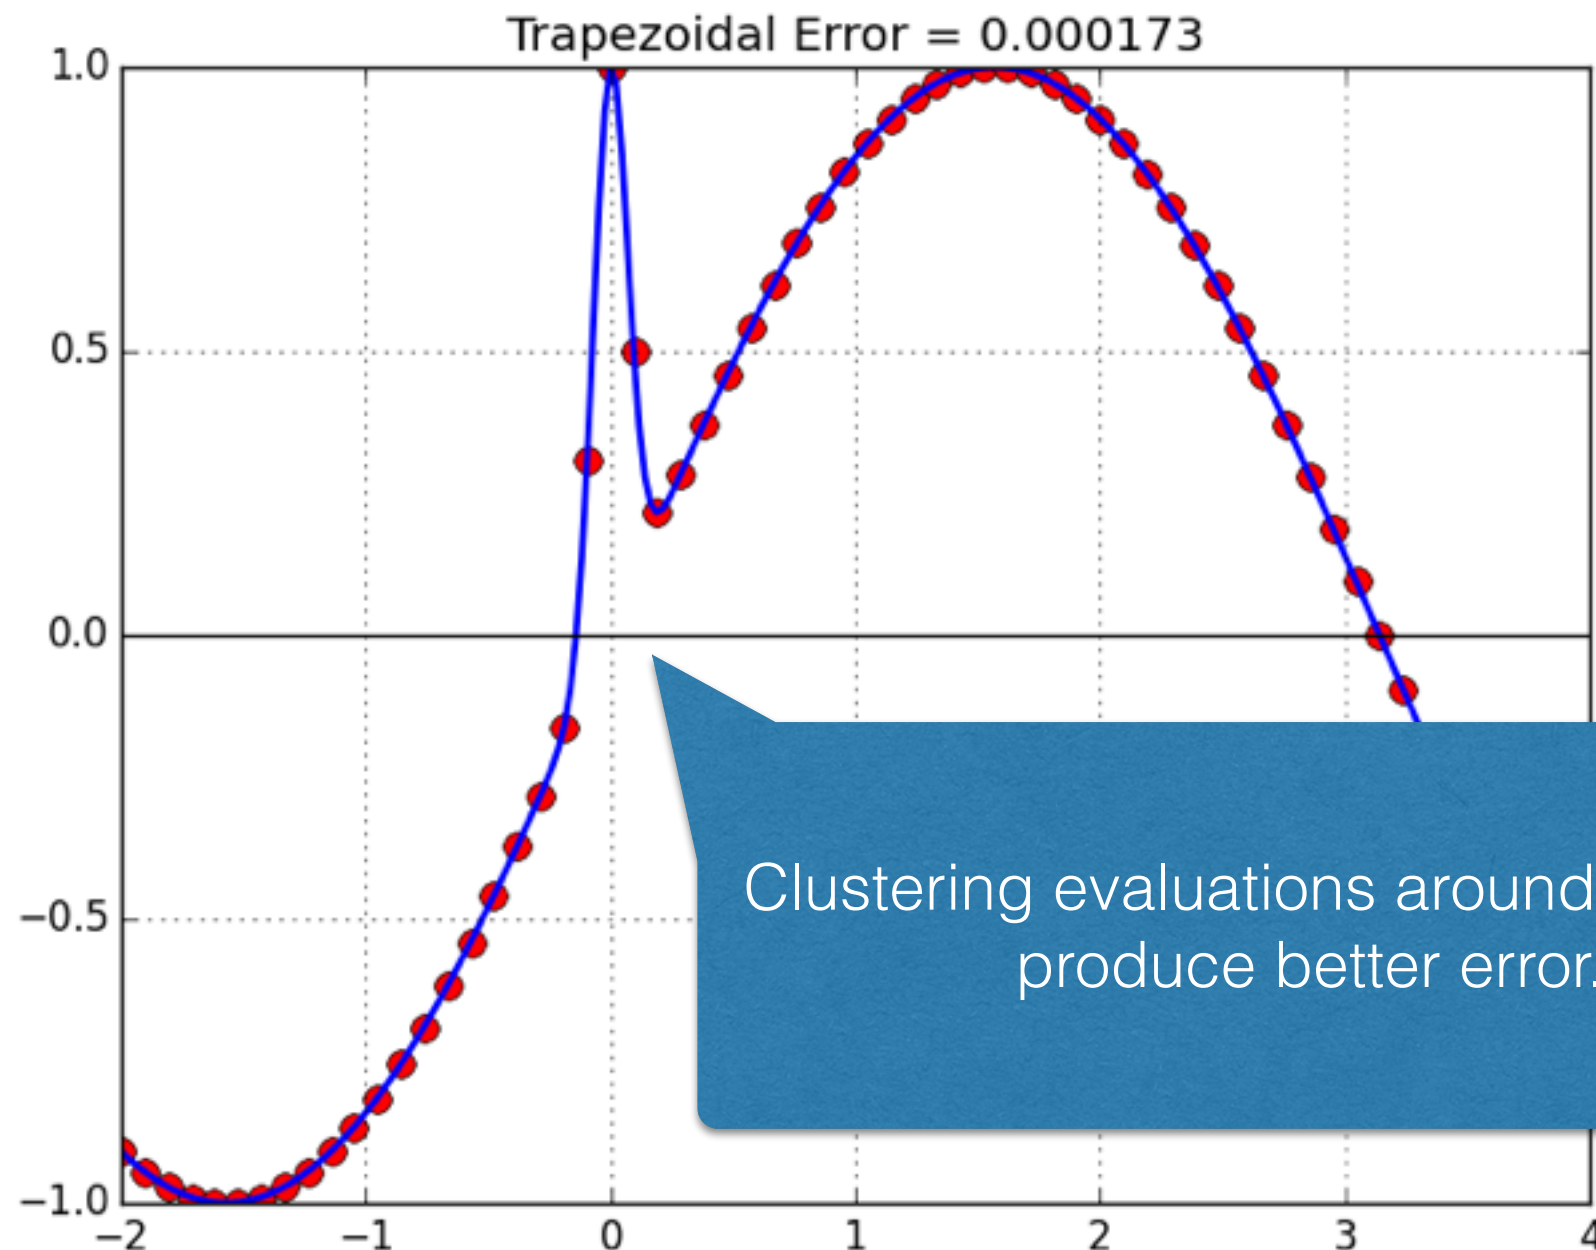
# Estimating Integral

- Trapezoidal rule:  $n=64$  points, equally spaced




# Estimating Integral

- Trapezoidal rule:  $n=64$  points, equally spaced



# Estimating Integral

- Trapezoidal rule:  $n=64$  points, equally spaced



# Brief Aside

$$\int_a^b f(x)dx \approx \text{trapz}(f, a, b)$$

- How to estimate approximation error?
- Compute both `trapz` and `simps` approx.
  - if abs. difference is large then **bad approx.**
  - if abs. difference is small then **good approx.**
- Reason:  $\text{error} \sim O(|b - a|^n)$

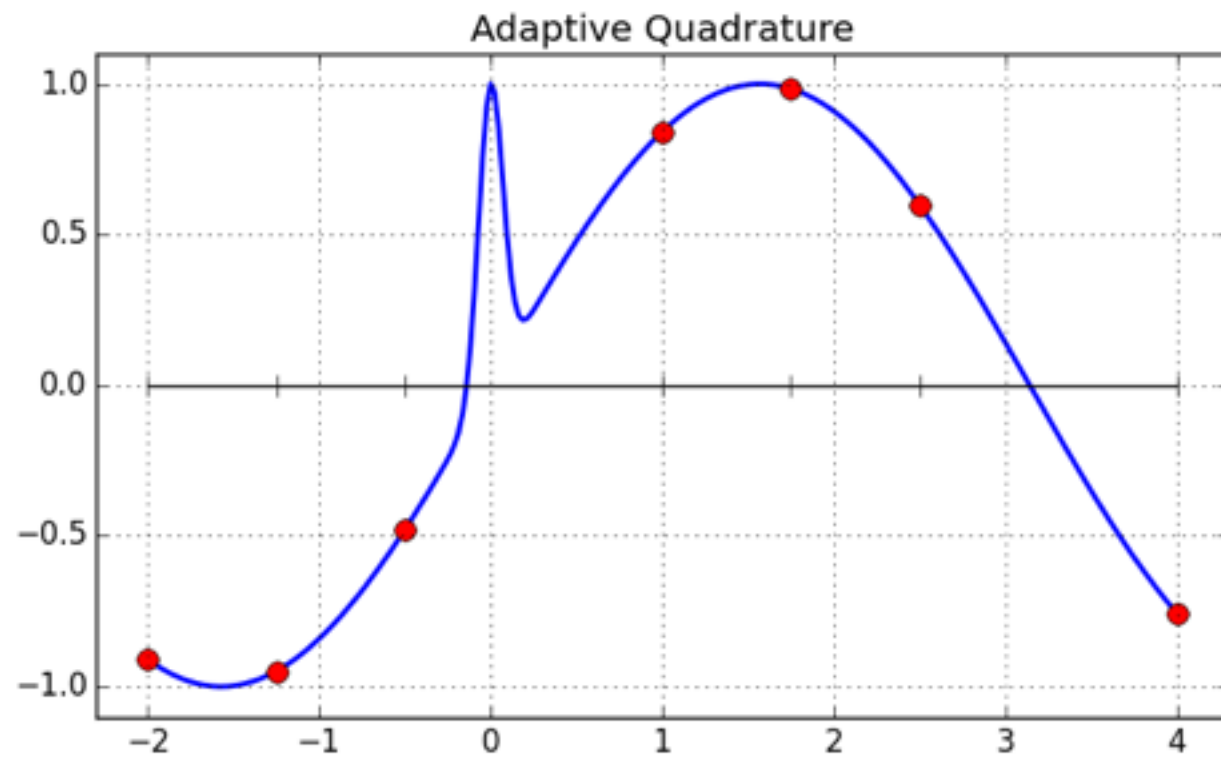
# Central Idea

- What if approx. error  $> \text{tol}$  ?


$$\int_a^b f(x)dx = \int_a^{(a+b)/2} f(x)dx + \int_{(a+b)/2}^b f(x)dx$$

- Idea: split interval in half:
  - suppose error in each half is less than  $\text{tol}/2$
  - then total error is less than  $\text{tol}$
- If  $\text{err} > \text{tol}/2$  then divide that half of the interval and repeat


# Quadrature tol=0.5




integral: 0.198245  
(actual: 0.414742)  
tolerance: 5.000000e-01  
est error: 2.223337e-01  
act error: 2.164973e-01



# Quadrature tol=0.1



integral: 0.407417  
(actual: 0.414742)  
tolerance: 1.000000e-01  
est error: 5.342968e-02  
act error: 7.325371e-03



# Demo

`adaptive_serial.c` and `test_serial.c`




# Parallelizing Using OpenMP

- **Naive Idea:** have each thread take a half of the domain
  - “domain decomposition pattern”
  - atomic / critical operations:
    - error calculation
    - integral summation

# Demo


`adaptive_serial.c` and `test_parallel.c`

# Par. Quad. tol=0.1




integral: 0.407417  
(actual: 0.414742)  
tolerance: 1.000000e-01  
est error: 5.342968e-02  
act error: 7.325371e-03


(Same as serial.)



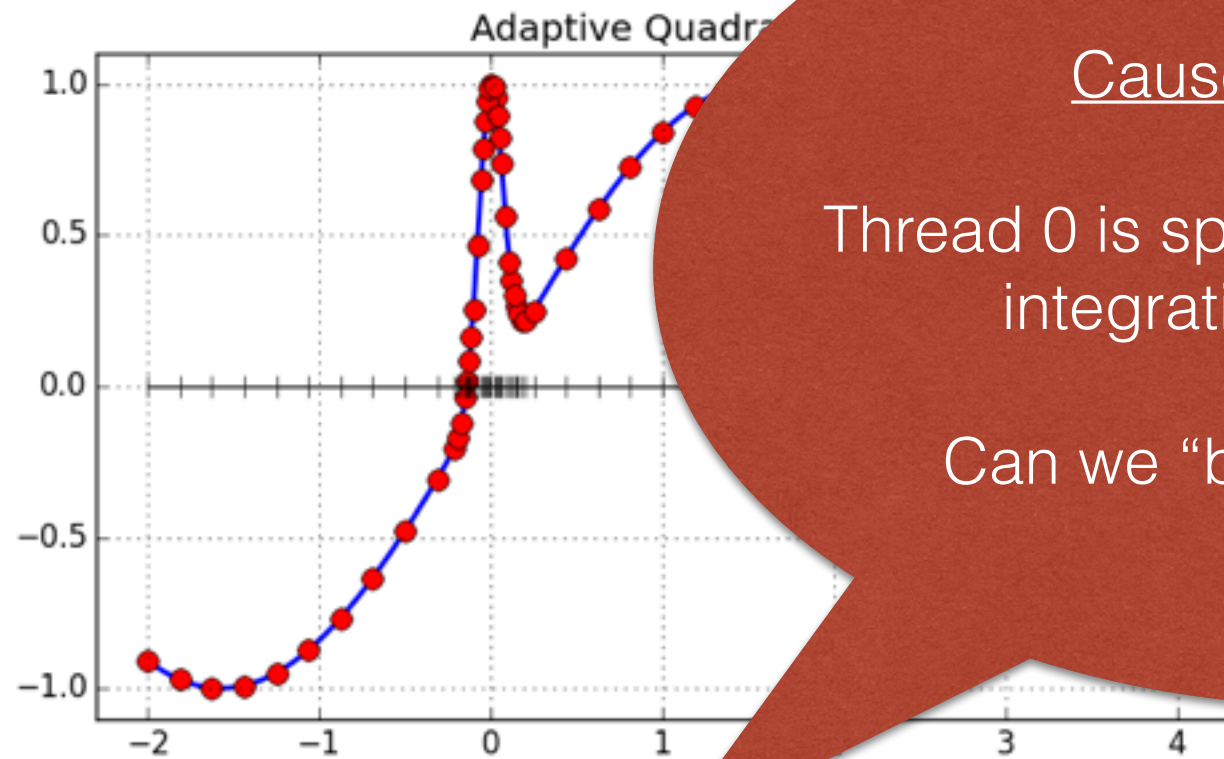
# Par. Quad. tol=0.02



integral: 0.414474  
(actual: 0.414742)  
tolerance: 2.000000e-02  
est error: 1.189837e-02  
act error: 2.678713e-04



# Par. Quad

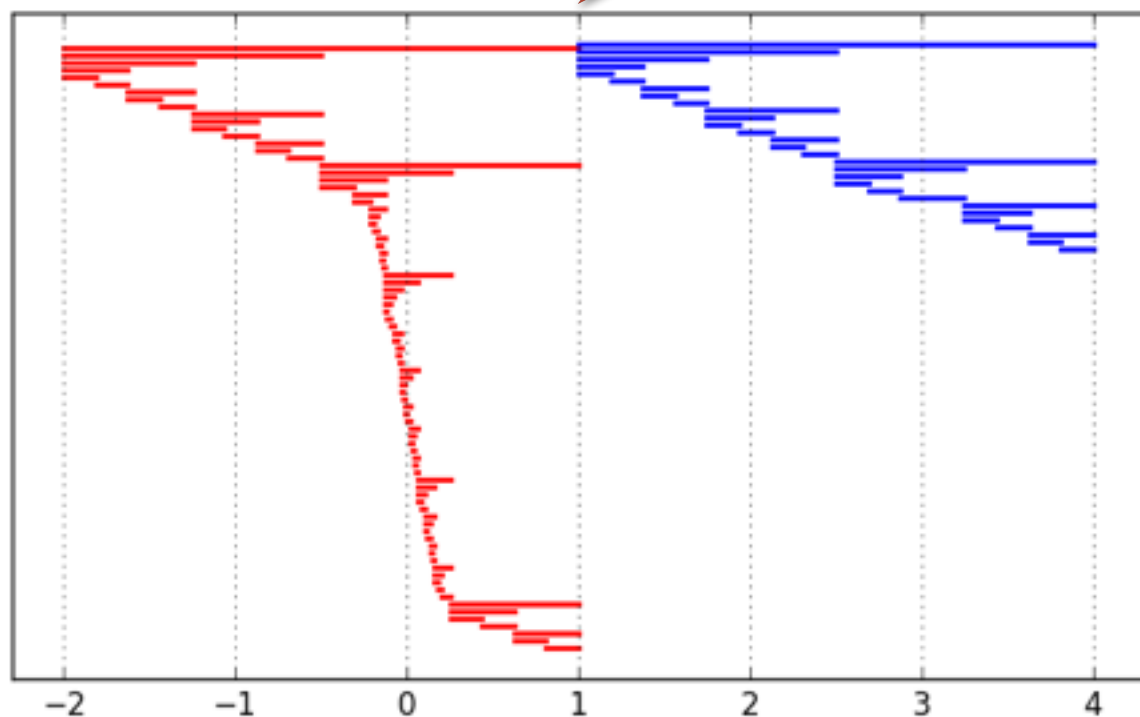


## Cause for Concern

Thread 0 is spending almost 3x time integrating as Thread 1.

Can we “balance” the load?

0.414474  
(actual: 0.414742)  
tolerance: 2.000000e-02  
est error: 1.189837e-02  
act error: 2.678713e-04



# OpenMP Sections

- Work-sharing Construct (like `omp for`) — each “section” is given to next available thread

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        // block of code
        #pragma omp section
        // block of code
    }
}
```

# OpenMP Sections

- Work-sharing Construct (like `omp for`) — each “section” is given to next available thread

```
#pragma omp parallel sections
```

```
{
```

```
    #pragma omp section
```

```
        // block of code
```

```
    #pragma omp section
```

```
        // block of code
```

```
}
```

Allowed to combine parallel and sections into one.

e.g.

```
#pragma omp parallel for
```

# OpenMP Sections

- What happens if `#threads != #sections`?
  - If more threads than sections?
  - If more sections than threads?




# OpenMP Sections

- What happens if `#threads != #sections`?
  - If more threads than sections?
    - Some threads will execute sections and some won't.
  - If more sections than threads?
    - OpenMP will allocate next available thread to next available section.


# Nested Parallelism

- Fork-Join Model



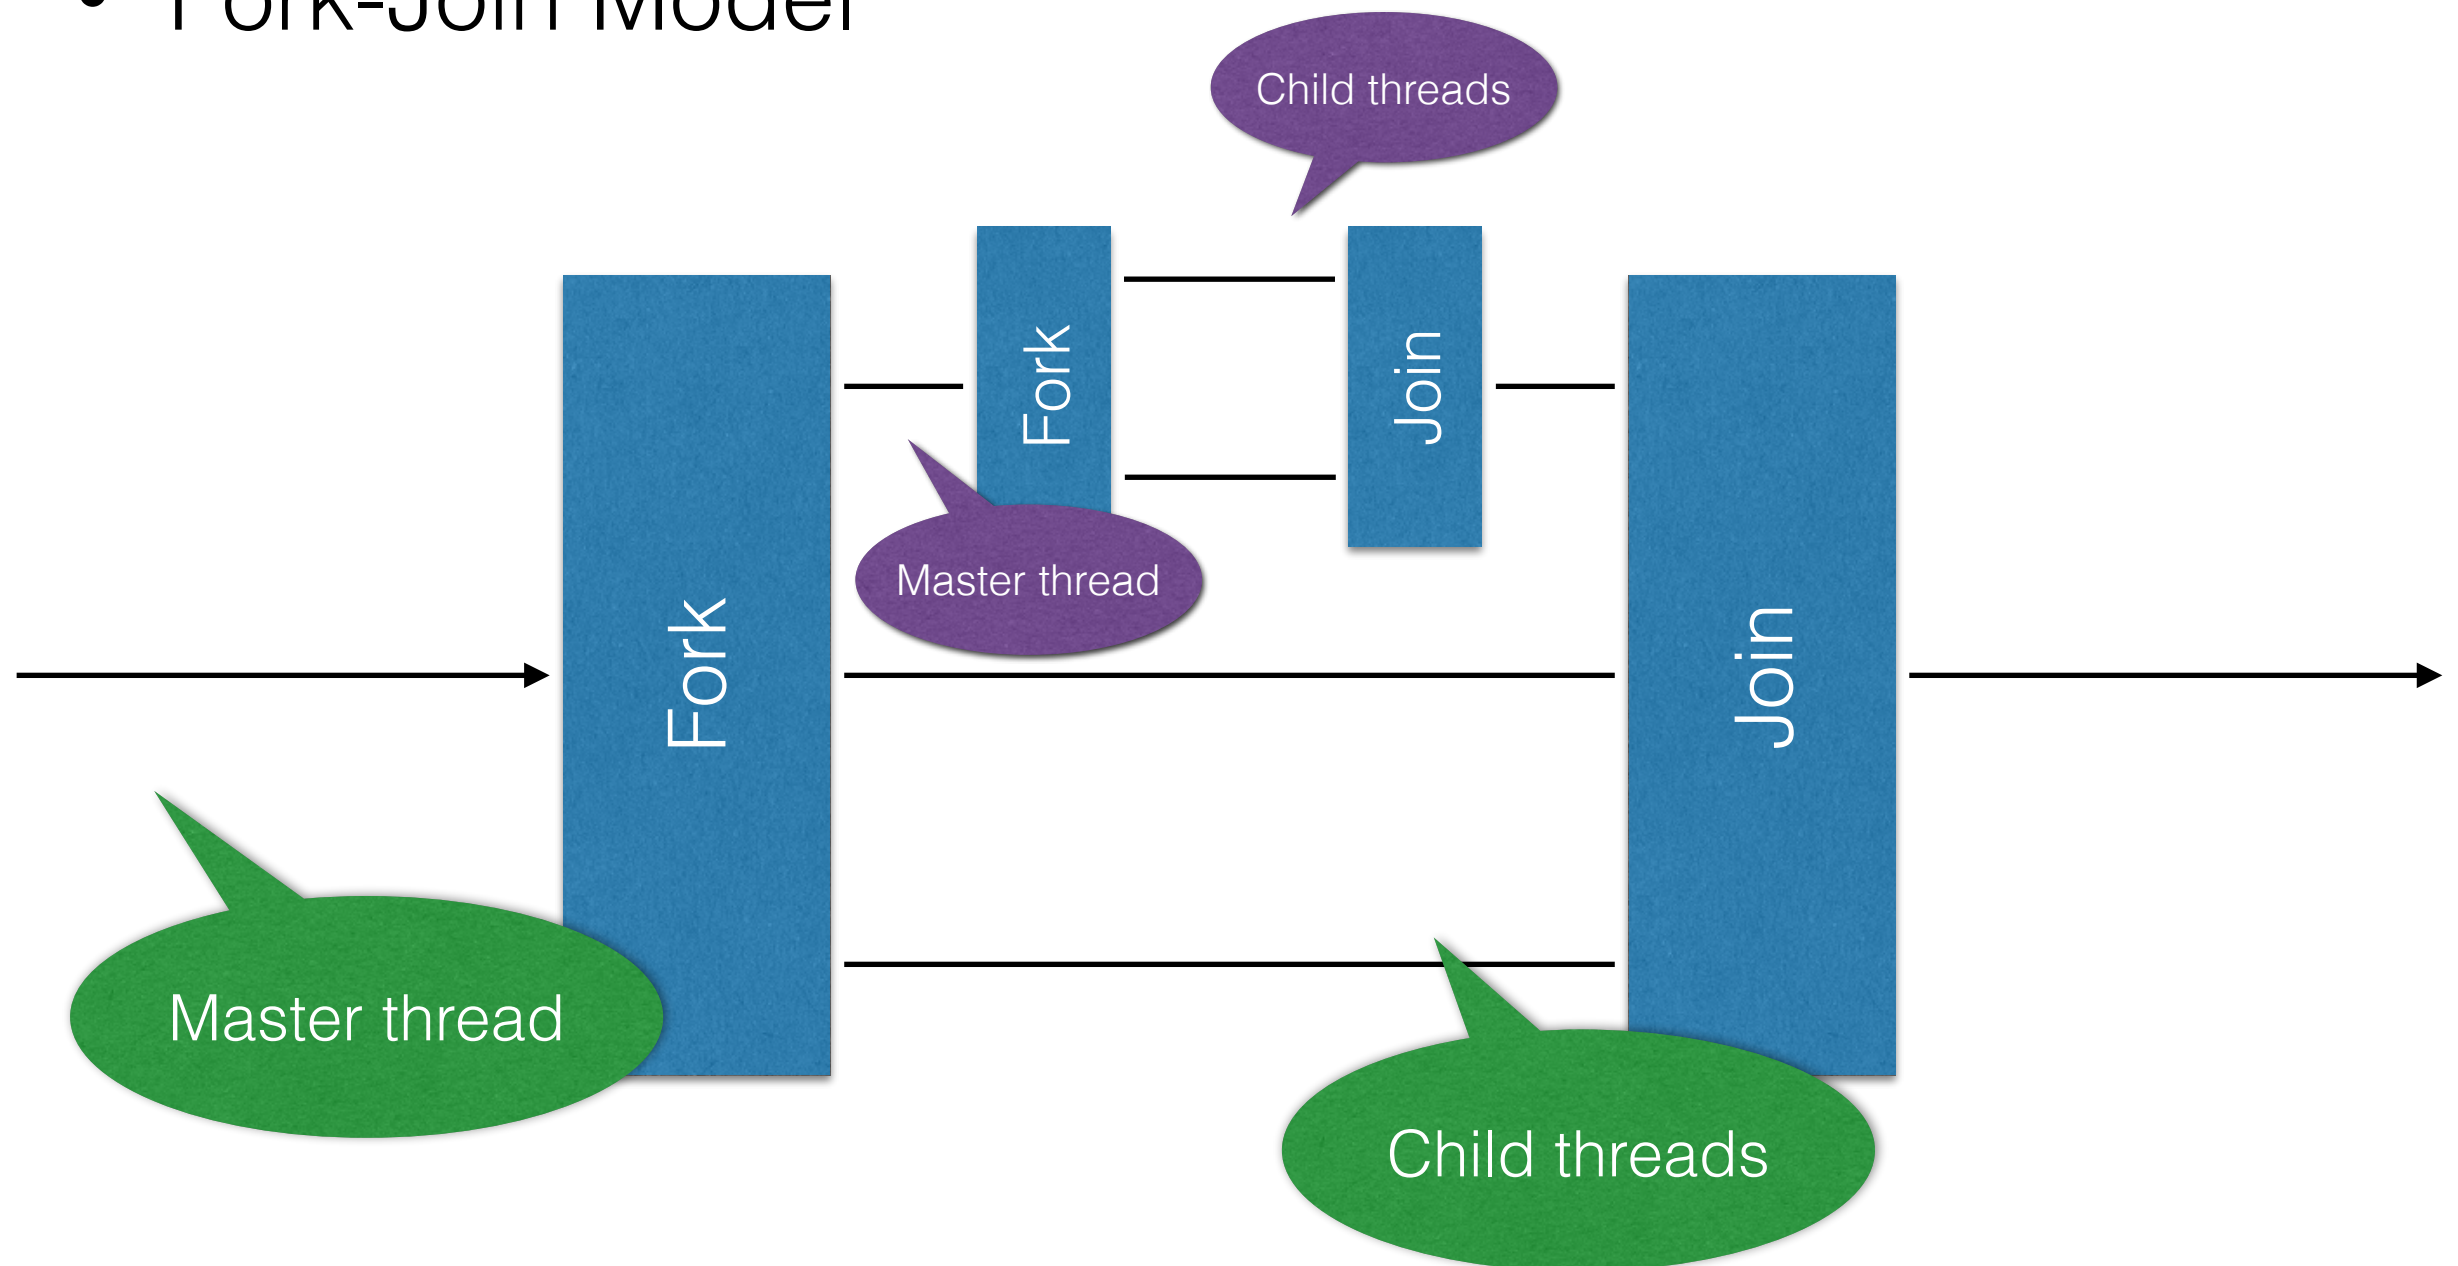
# Nested Parallelism

- Fork-Join Model



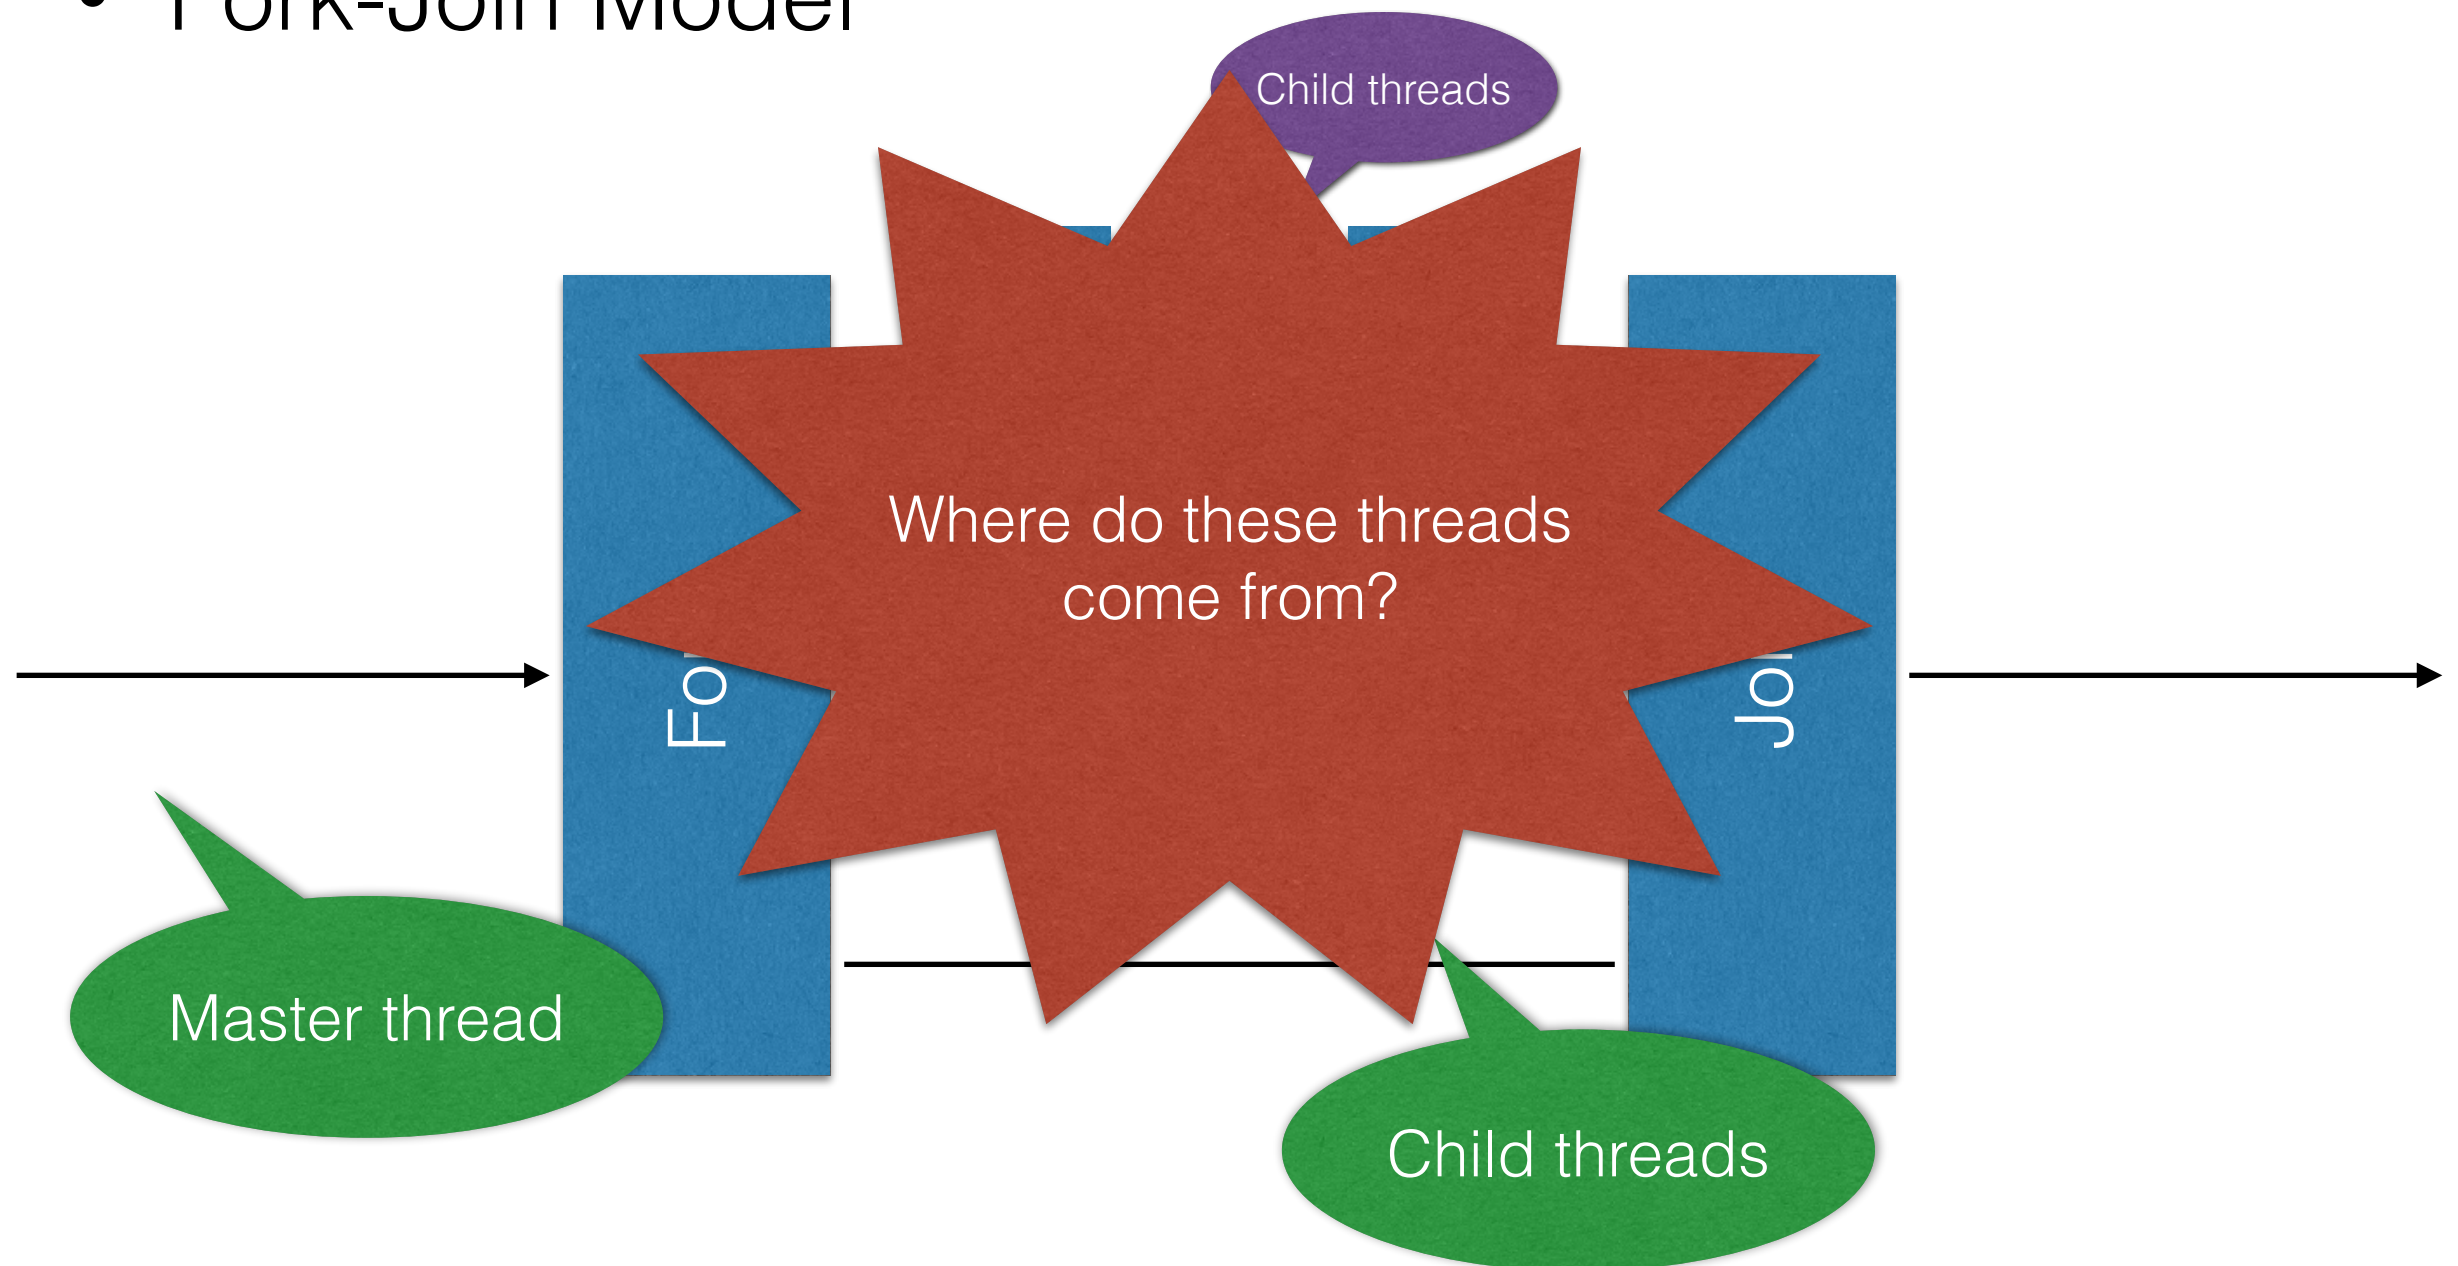
# Nested Parallelism

- Fork-Join Model



# Nested Parallelism

- Fork-Join Model



# Nested Parallelism

- OpenMP maintains a “*pool*” of threads
  - encounter `omp parallel` —> take idle threads
  - encounter another `omp parallel` —> take idle threads
- After JOIN, threads are returned to pool.

# Nested Parallelism



- Possible to create many total threads!
- Many ways to deal with this, see online resources.

# Demo

`adaptive_parallel.c` and `test_balanced.c`





# Balanced Quad. tol=0.1



```
num_threads = 2  
integral:    0.407417  
(actual:    0.414742)  
tolerance:  1.000000e-01  
est error:  5.342968e-02  
act error:  7.325371e-03
```

# Balanced Quad. tol=0.1





num  
in  
(a  
to  
es

Work is shared evenly between  
Thread 0 and Thread 1.

(Each nested omp section  
given to next available thread.)

# Balanced Quad. tol=0.01



```
num_threads = 2  
integral:    0.414743  
(actual:    0.414742)  
tolerance:  1.000000e-02  
est error:  4.902247e-03  
act error:  9.255080e-07
```

# Thoughts on MPI

- OpenMP has built in tools for thread management
- Possible with MPI via `MPI_Comm_spawn()` and master —> child design pattern
- Many recursive functions can be re-implemented w/o recursion using a manual stack.
  - parallel producer — consumer model

# Thoughts on MPI

