


Lecture #18 - GPU Programming

AMath 483/583

GPUs

- Well-suited for SIMD
- Massively parallel
 - many “cores”




NVIDIA Tesla P100: *3584 cores*

- specializes in latency hiding
- Highly energy efficient (Gflops / \$ / Watt)


CUDA

- NVIDIA's superset of C/C++
- Write CPU and GPU code in same source file
- Creates “Host Program” on CPU and “Device Program” on GPU.
- communication between host and device similar to MPI



NVIDIA Tesla C2075


- 448 cores
- 1150 MHz clock
- 144 GB/sec mem. bandwidth
- 6 GB RAM
- 515 GFLOPS double precision



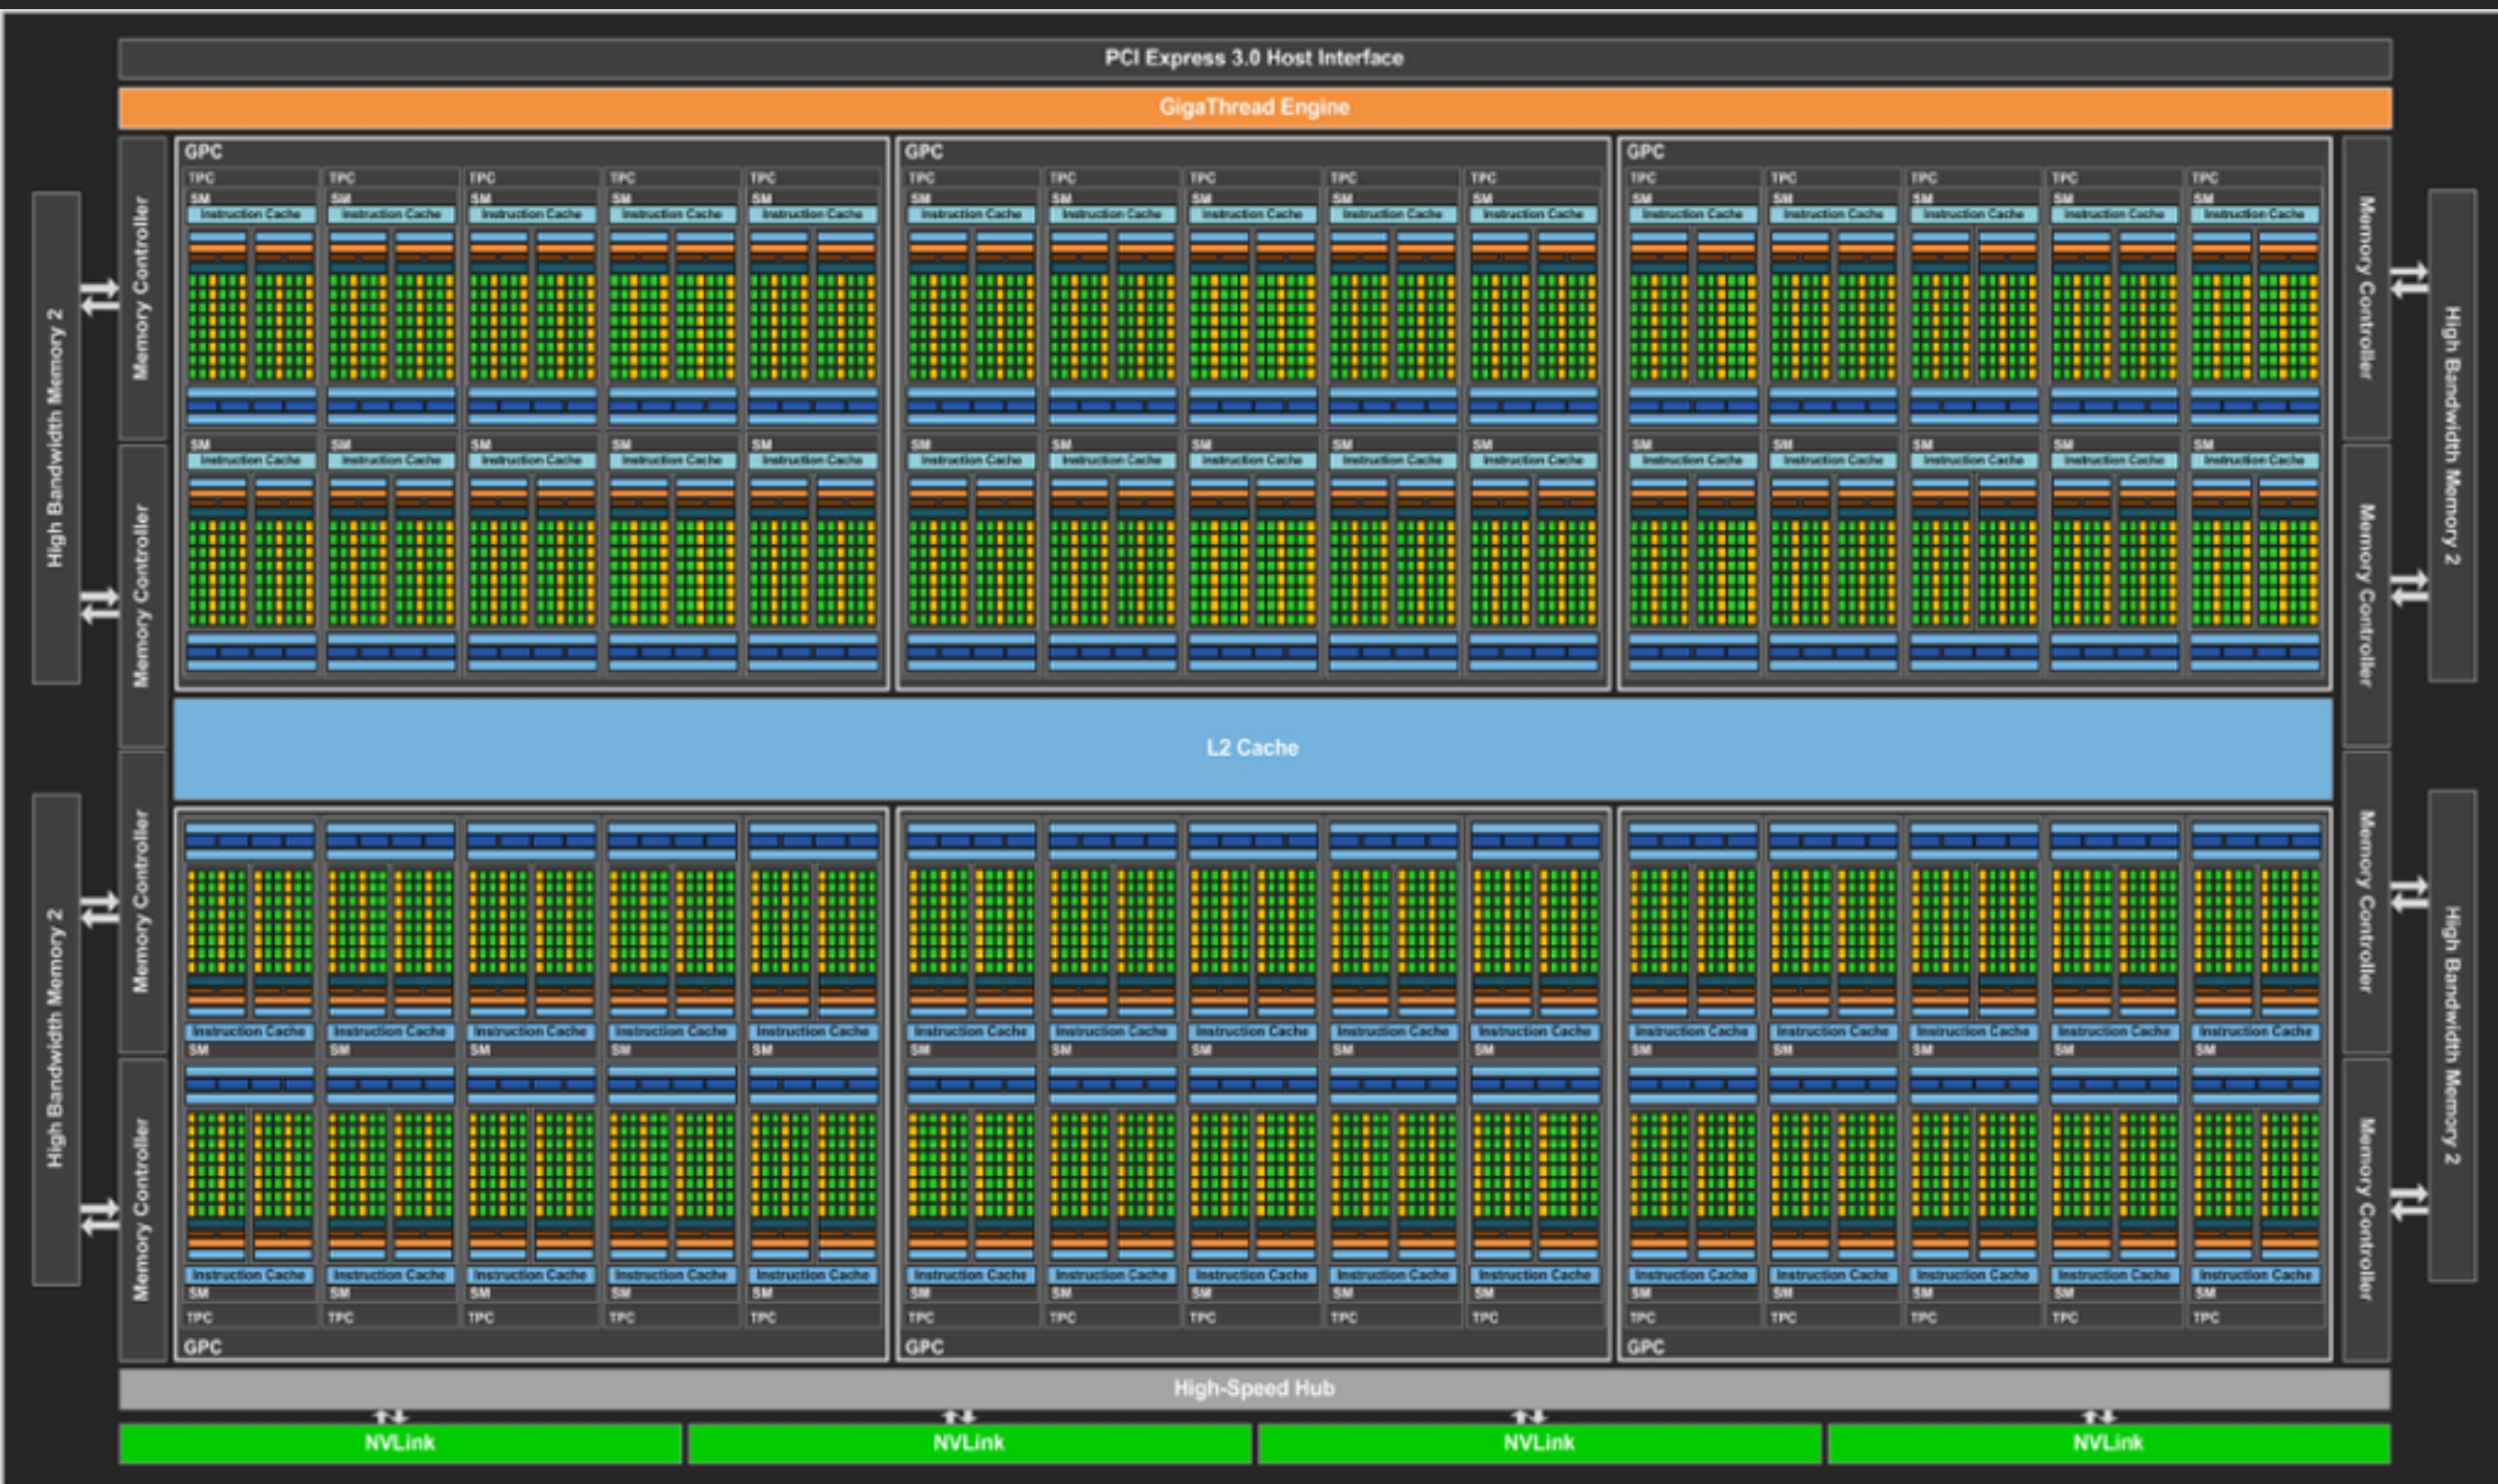
Used in this week's demos.
x4 on americano.amath

NVIDIA Tesla P100

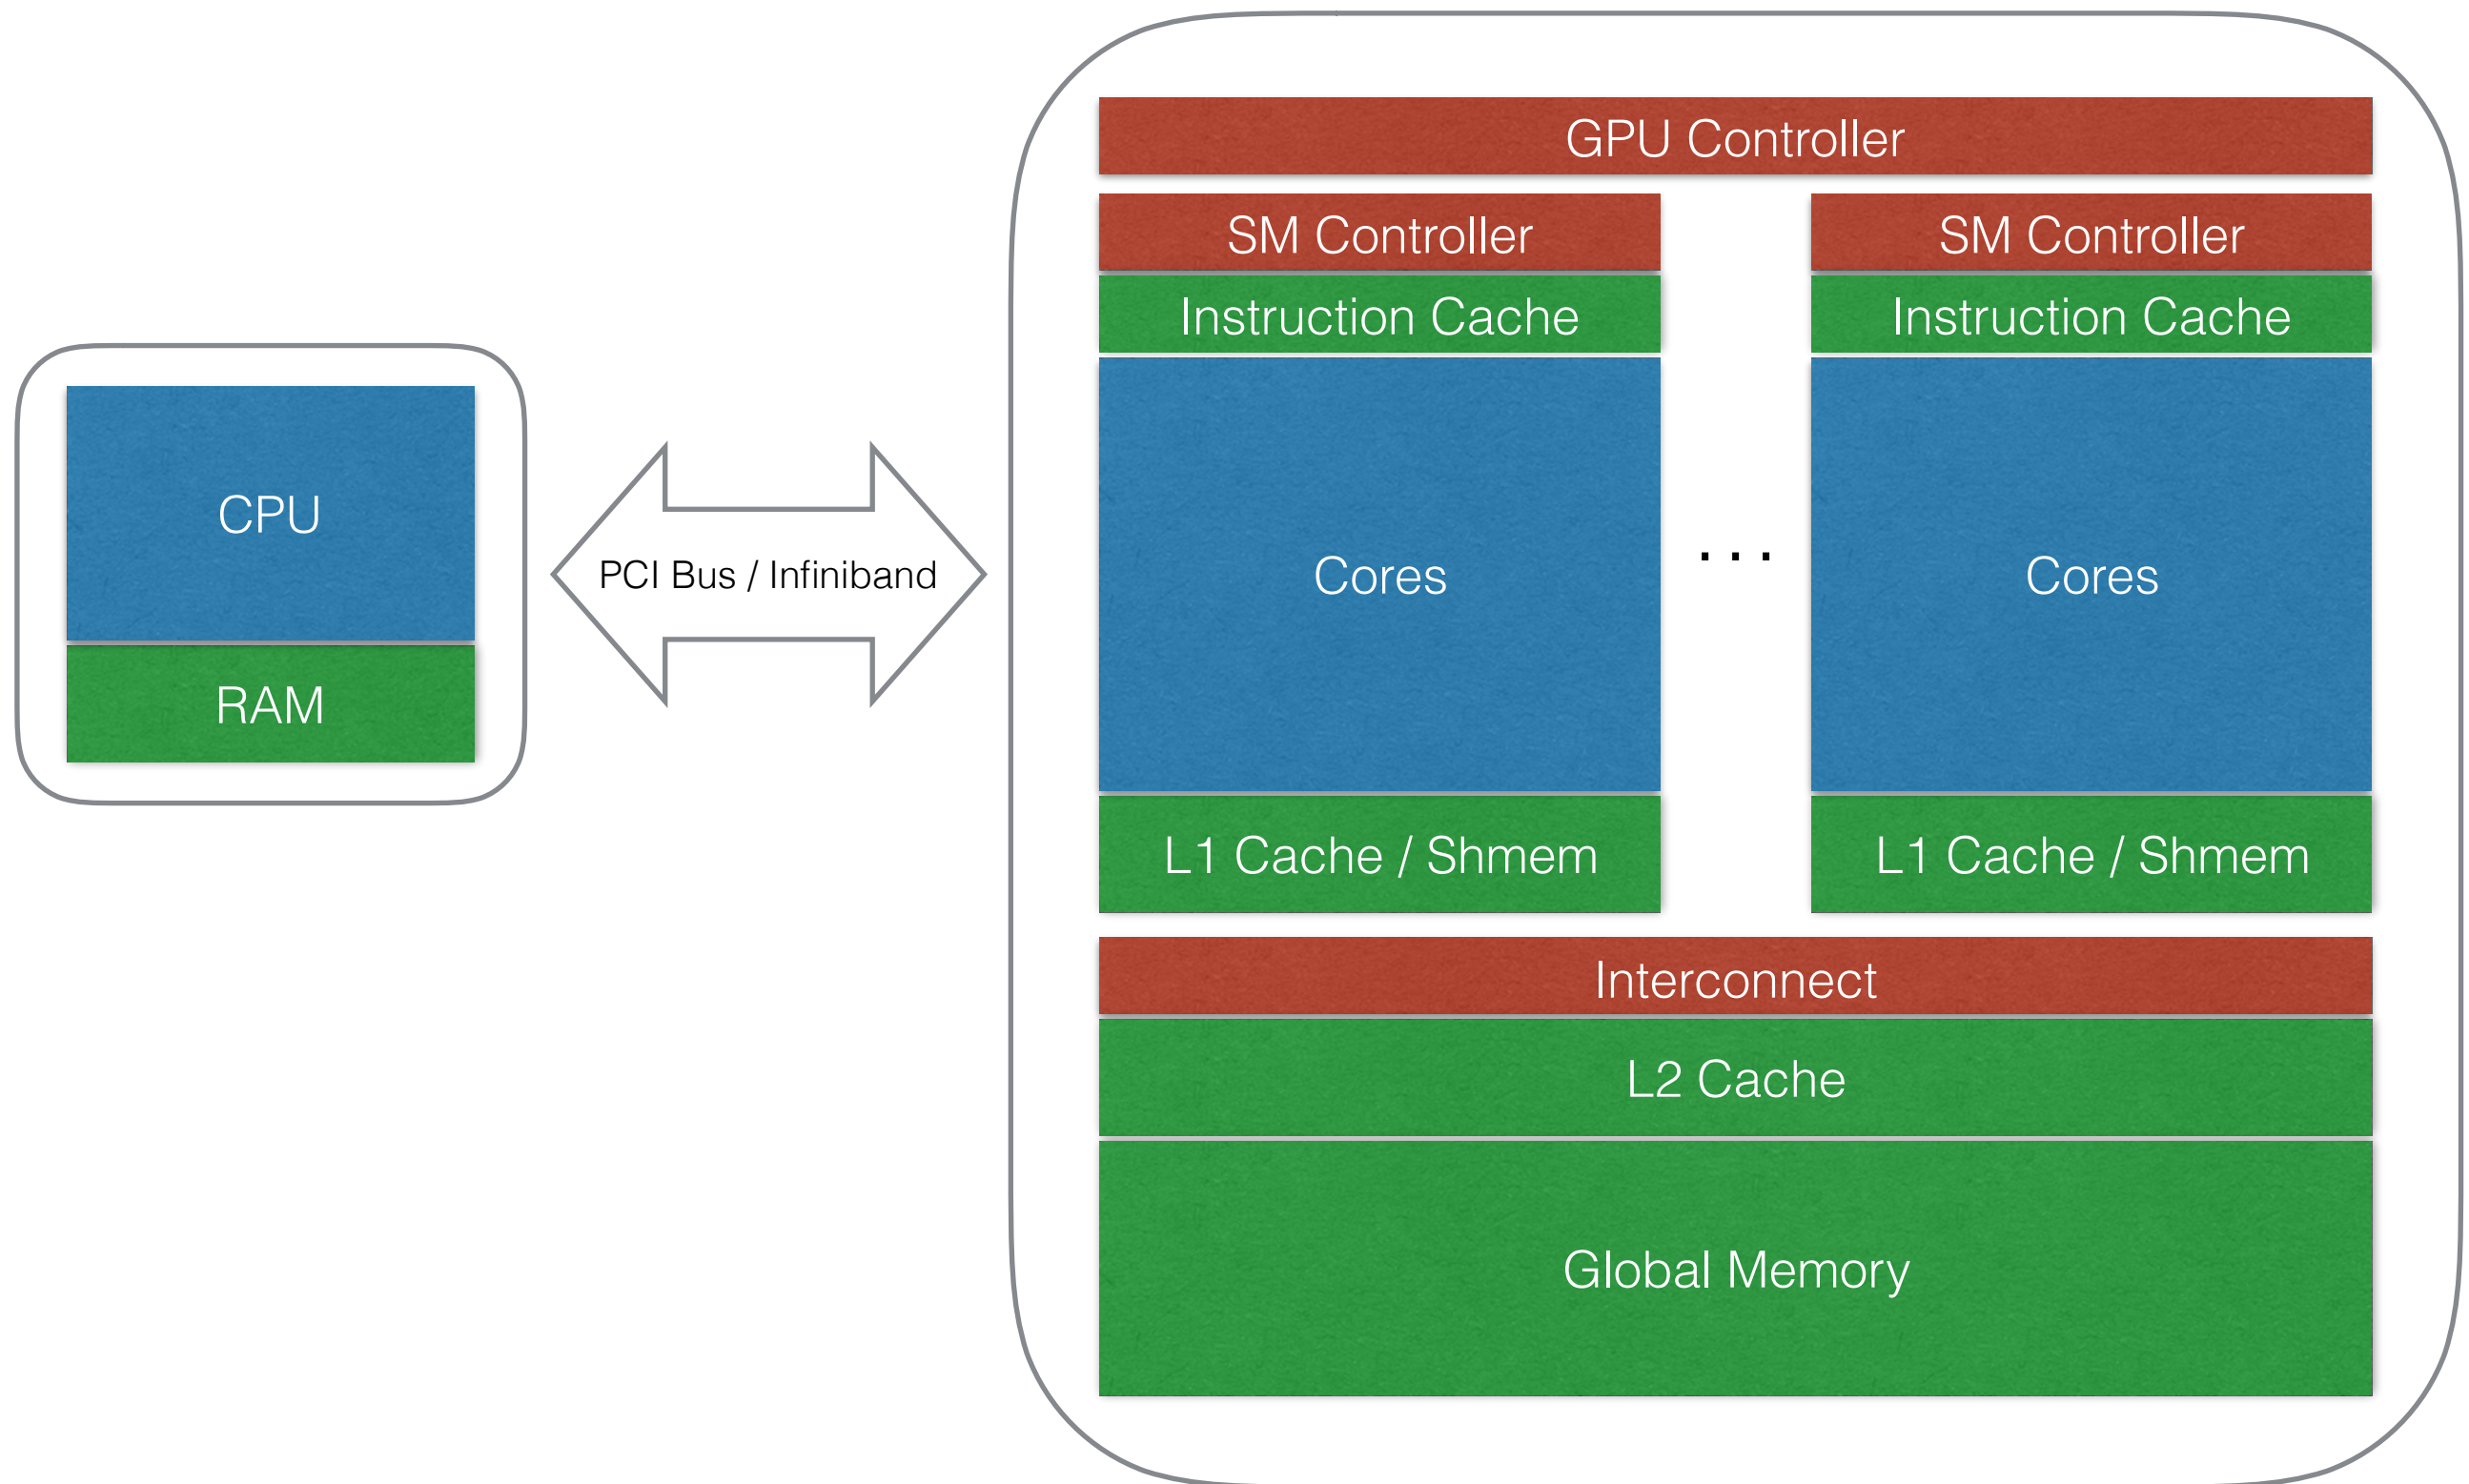
- 3584 cores
- 1328 MHz clock
- 720 GB/sec mem. bandwidth
- 16 GB RAM
- 5.3 TFLOPS double precision



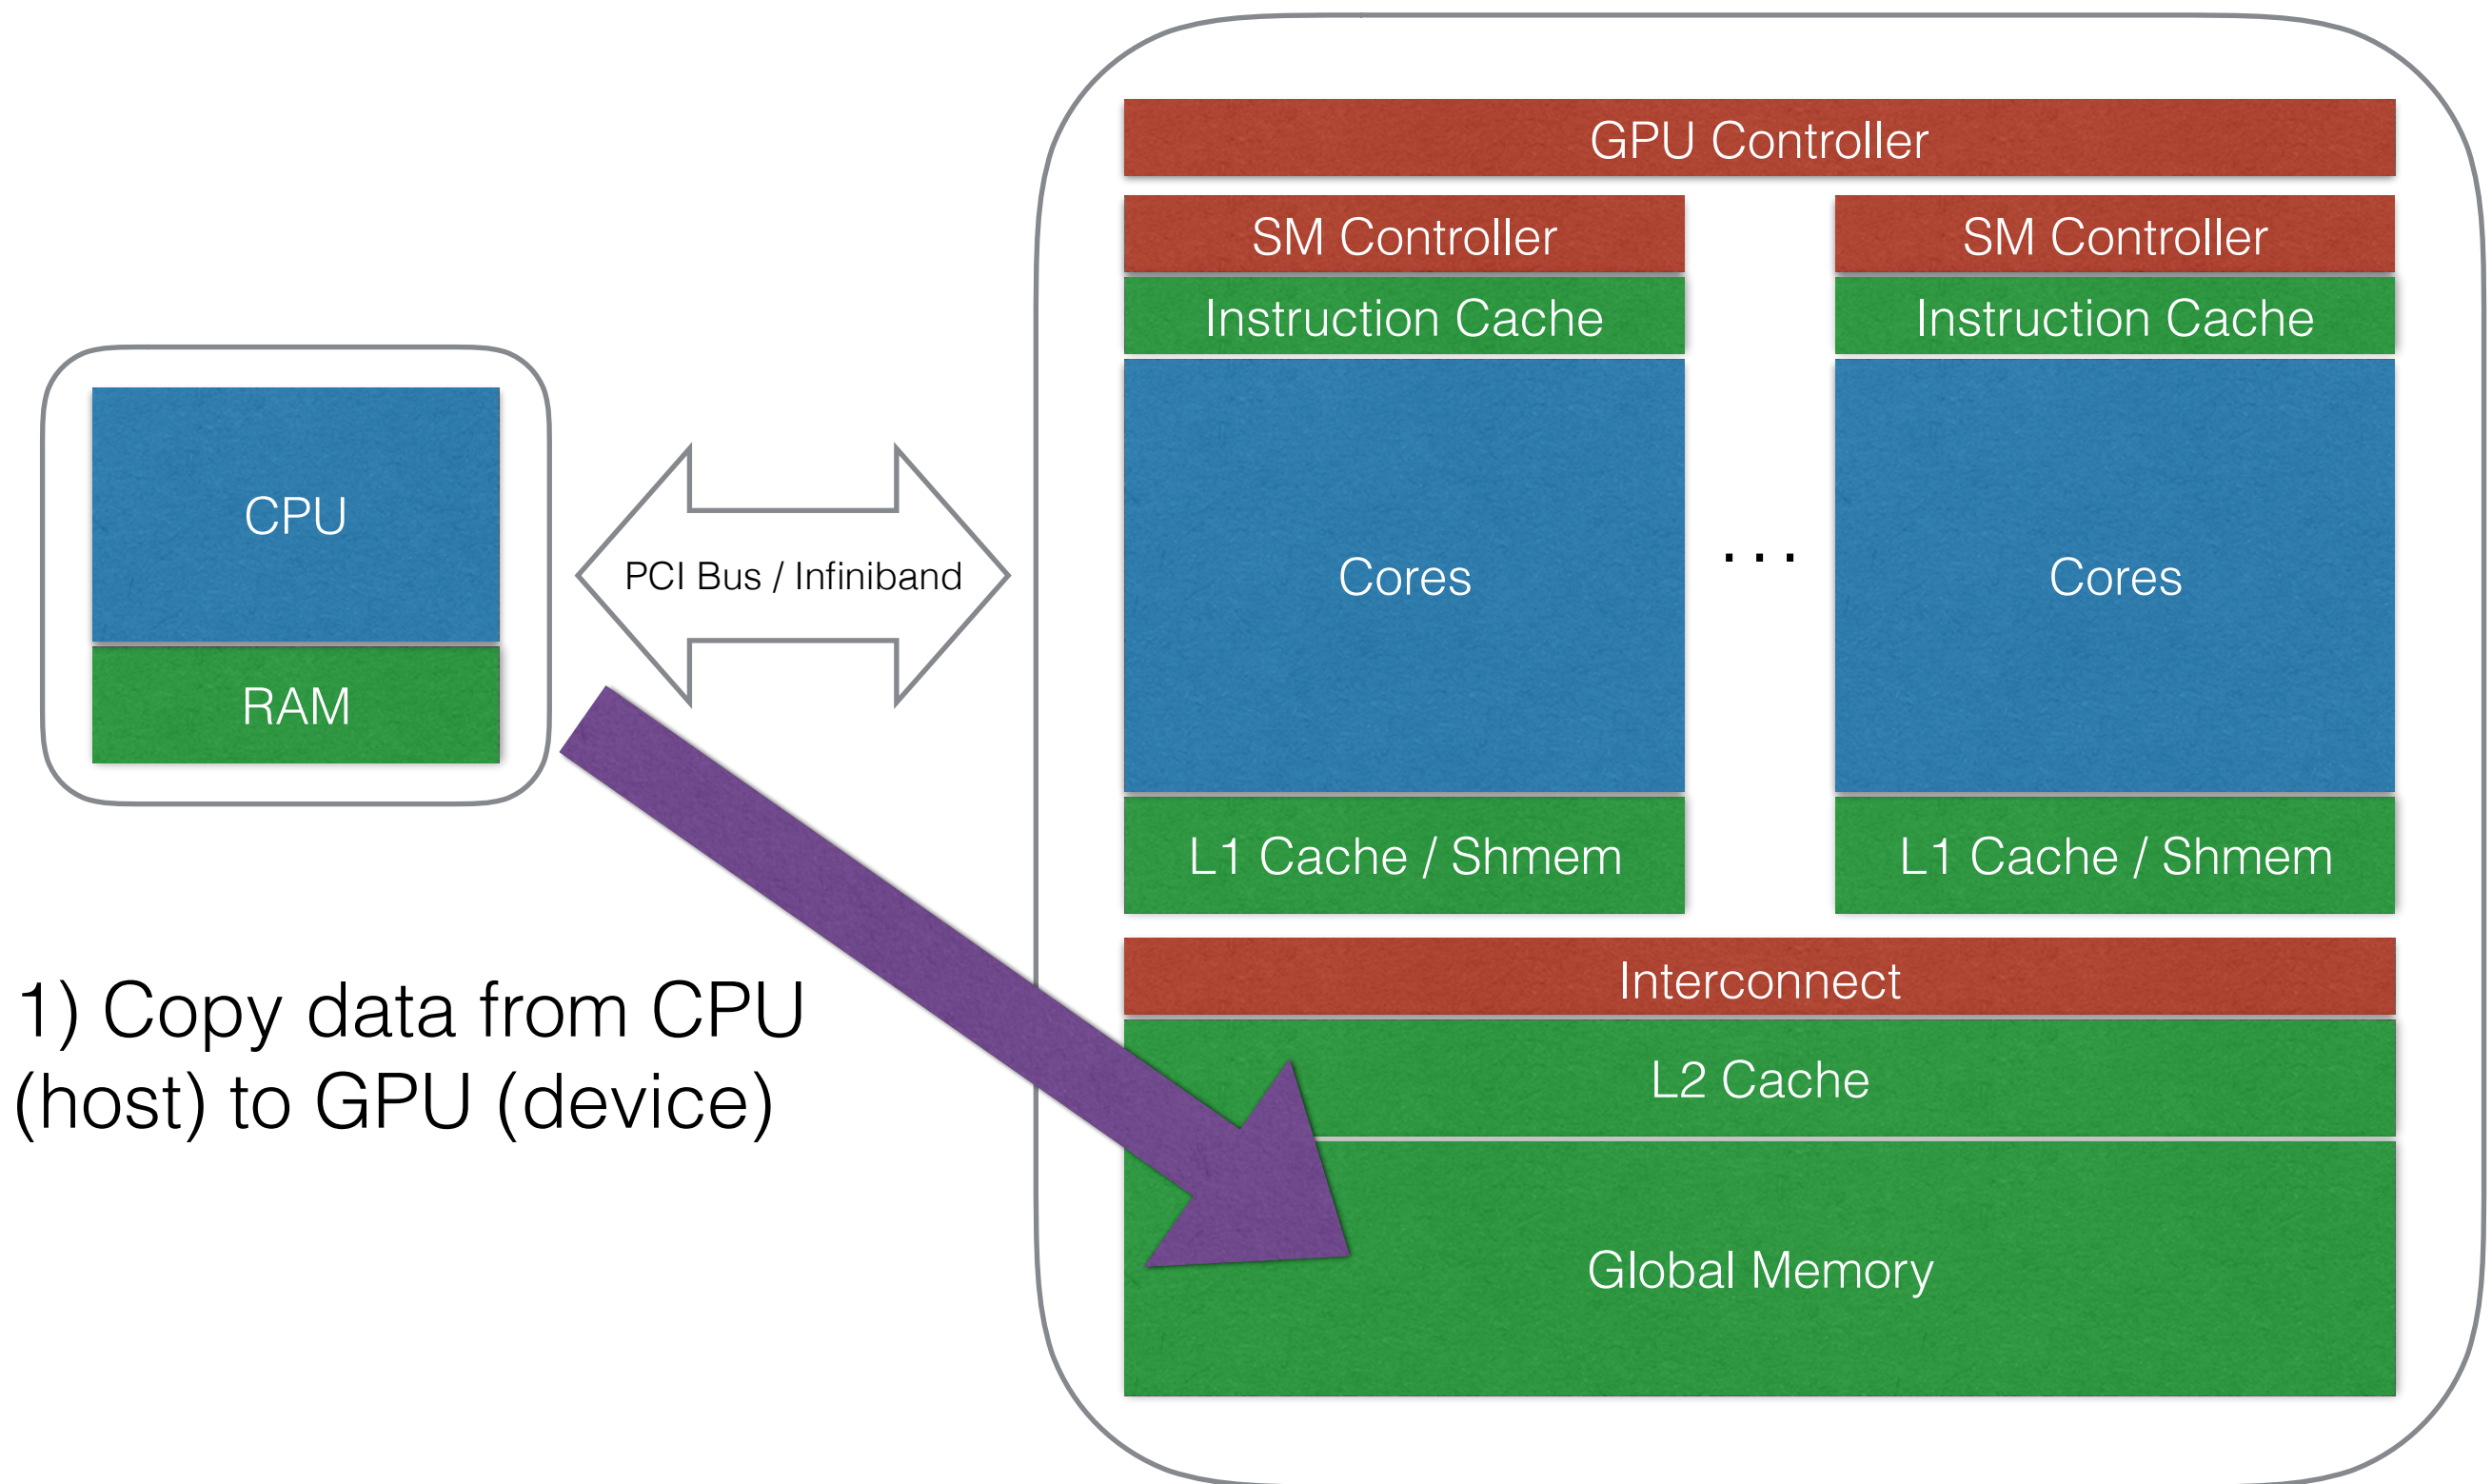
Computing power of top
supercomputer in 2000



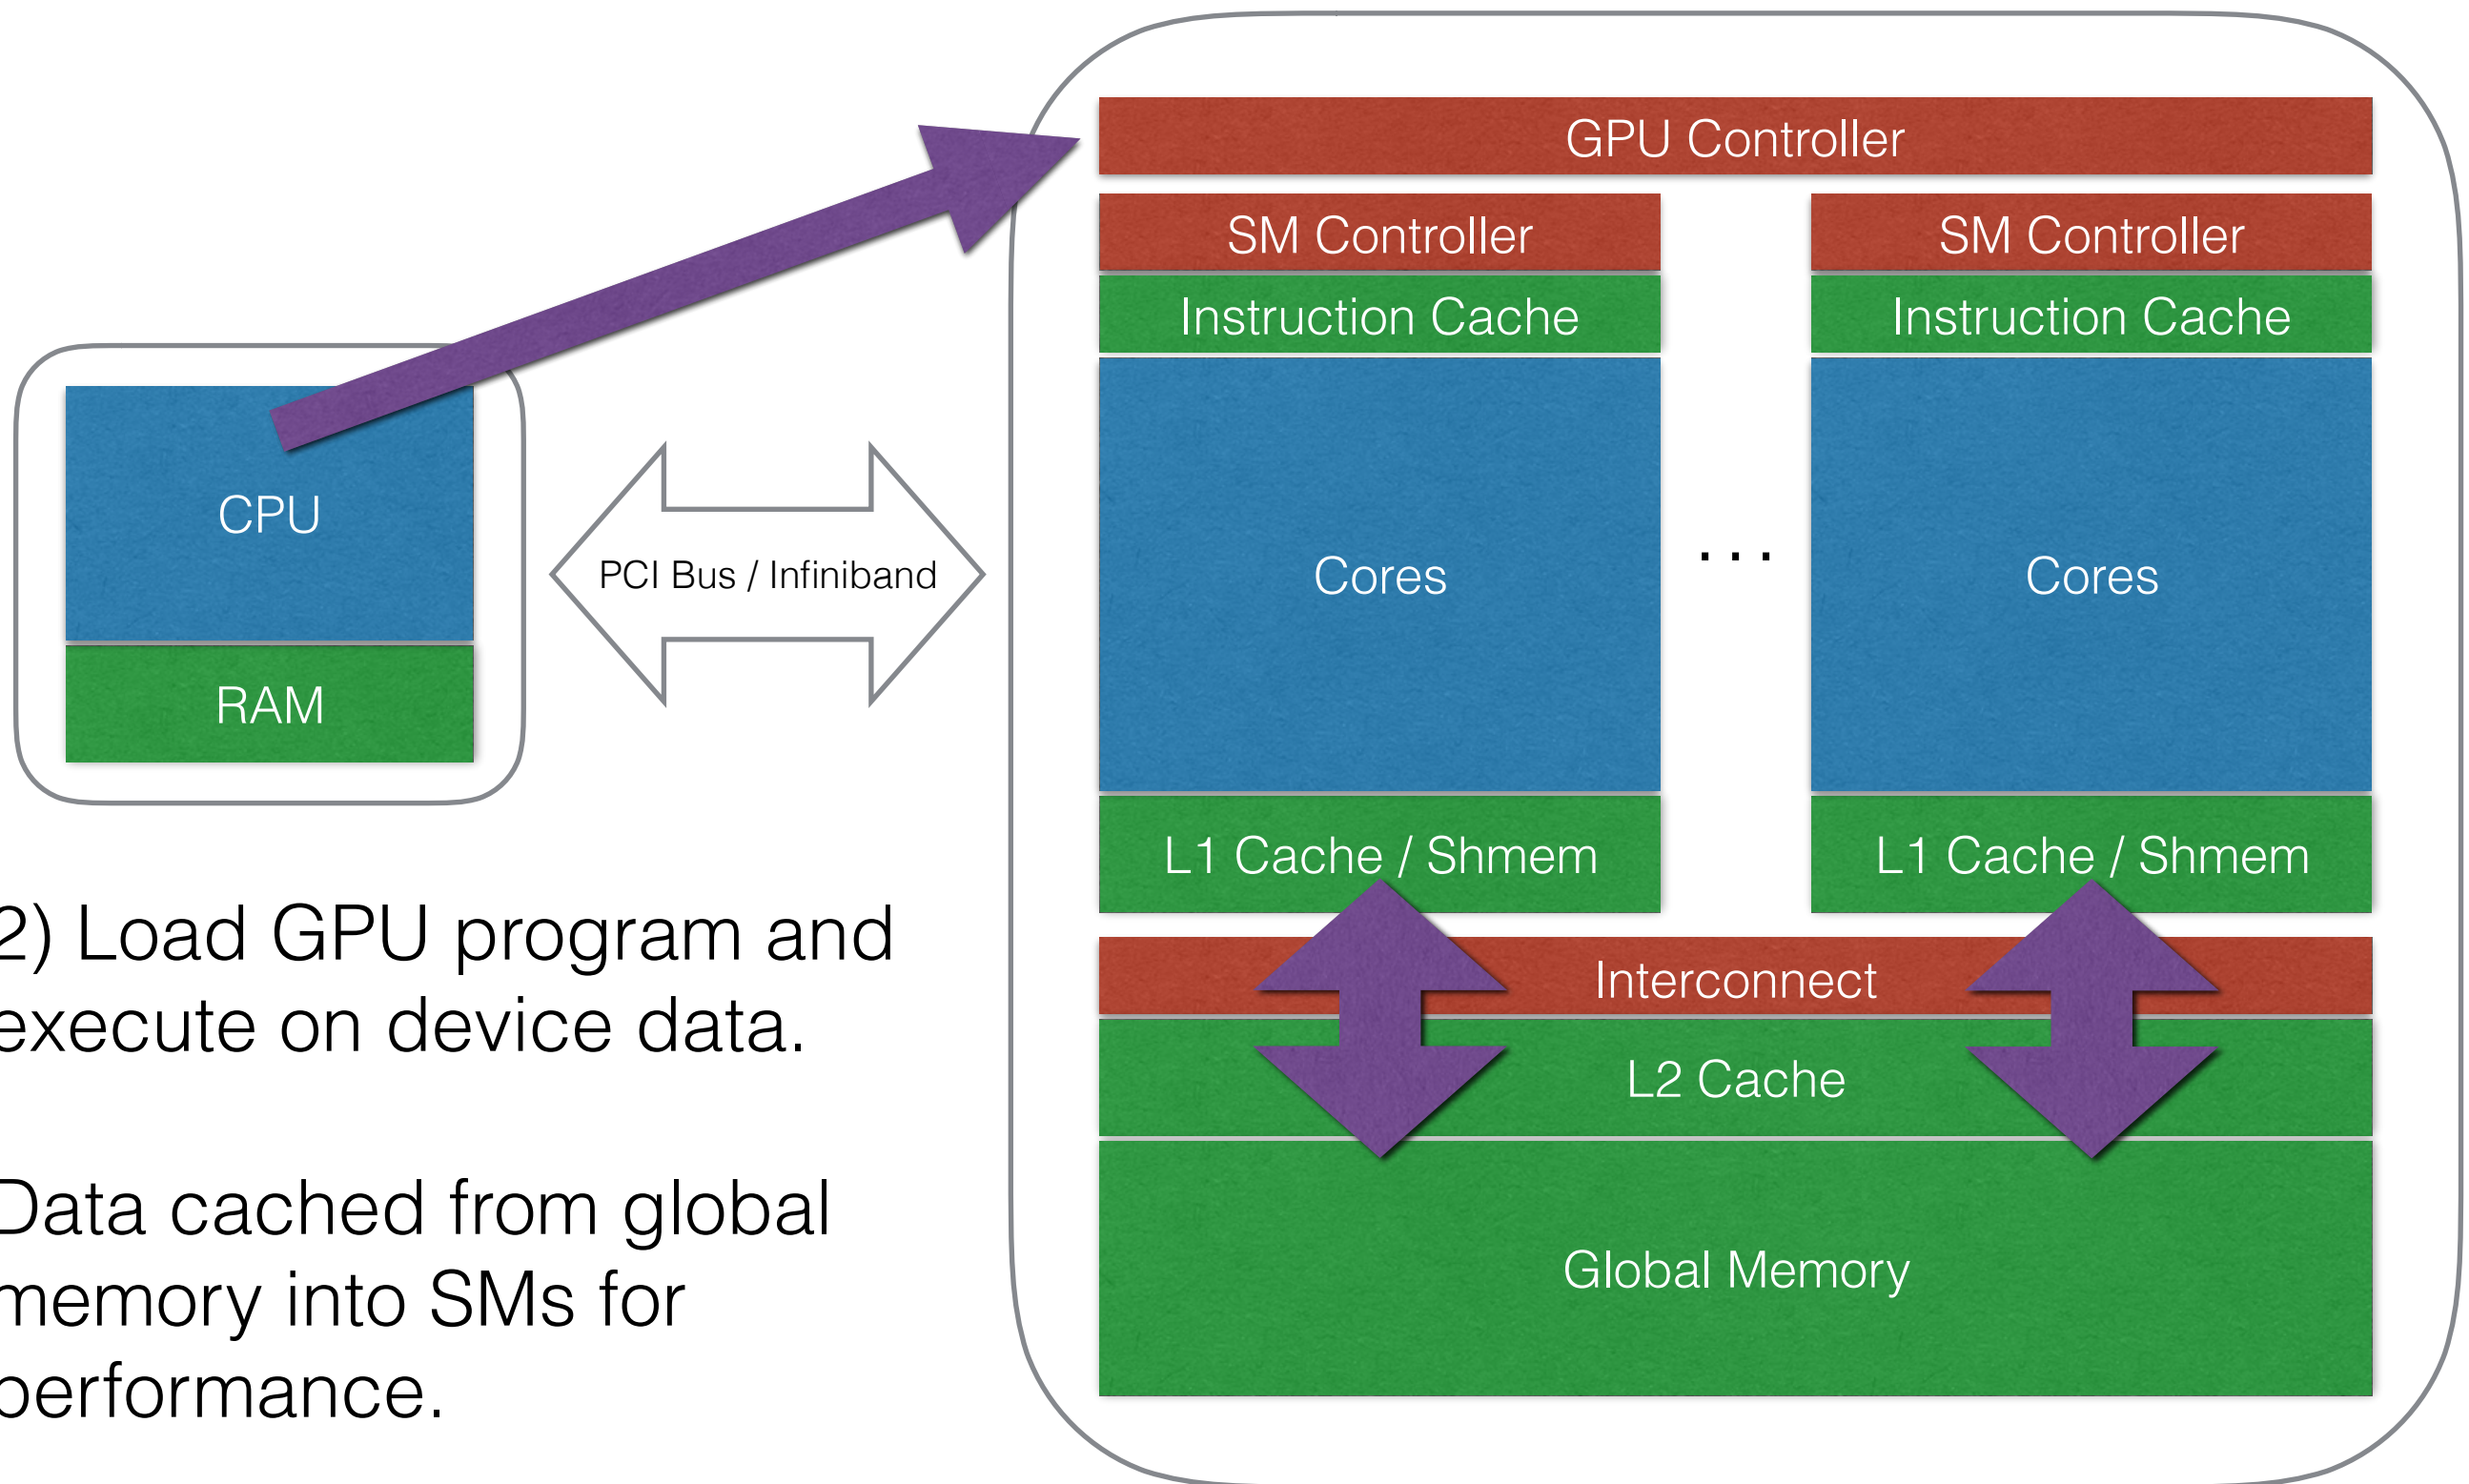
CPU vs. GPU



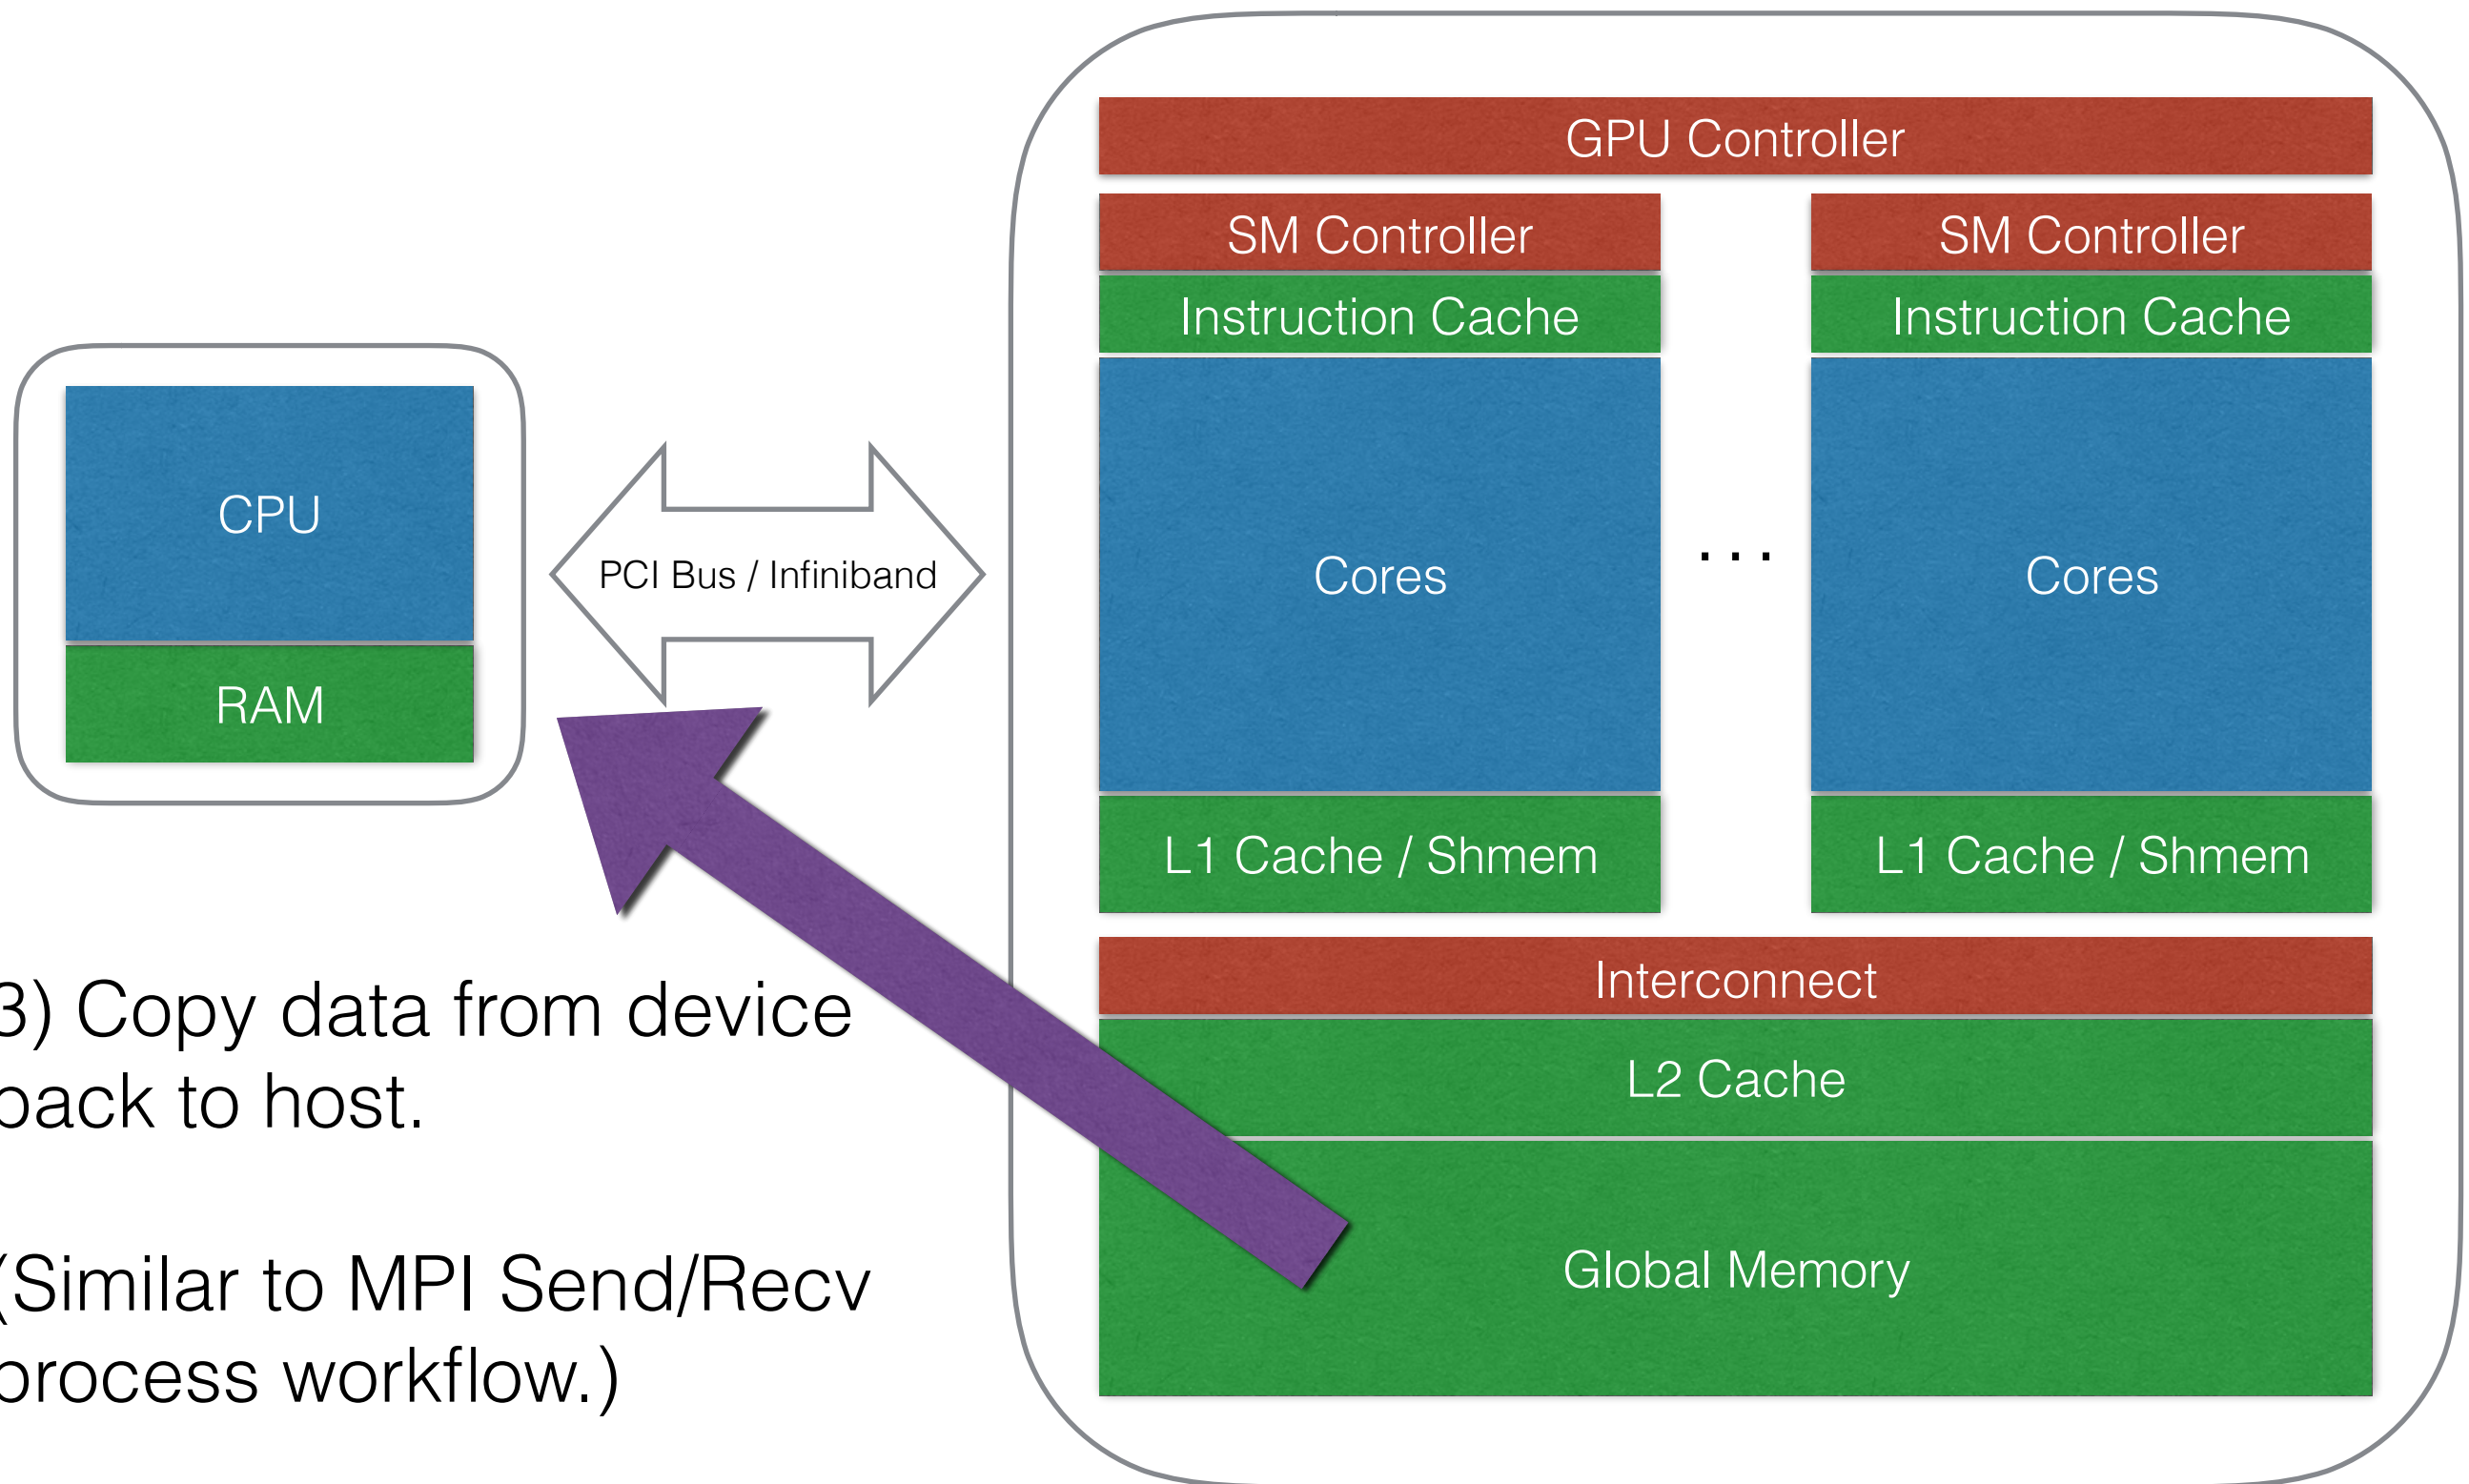
Simple Processing Flow



Simple Processing Flow



Simple Processing Flow



Terminology

- “Host-side” — code running on CPU. Directs GPU calls and memory transfer.
- “Device-side” — code running on GPU. Only interacts with device-side memory.
- “Kernel” — a function that runs on GPU. Invoked by host-side call. (Sometimes, device-side call.)

Hello World

```
#include <cuda.h>
```

95% of host-side functions defined here

```
__global__ void cuda_hello(void) {  
    // print a character buffer from the GPU!  
    printf("Hello, world!");  
}
```

```
int main(void) {  
    // call the CUDA kernel from the GPU  
    cuda_hello<<<1,1>>>();  
  
    // wait for all CUDA kernels to finish  
    cudaDeviceSynchronize();  
    return 0;  
}
```


Hello World

```
#include <cuda.h>
```

```
__global__ void cuda_hello(void) {  
    // print a character buffer from the GPU!  
    printf("Hello, world!");  
}
```

Aside: printing from GPU is bananas!

```
int main(void) {  
    // call the CUDA kernel from the CPU  
    cuda_hello<<<1,1>>>>();  
  
    // wait for all CUDA kernels to finish  
    cudaDeviceSynchronize();  
    return 0;  
}
```

A CUDA “kernel”

__global__ indicates that this function is run on the GPU

Hello World

```
#include <cuda.h>
```

```
__global__ void cuda_hello()  
// print a character  
printf("Hello, world");  
}
```

Invocation of the Kernel

For now, the <<<1,1>>> just means to run the kernel on a single GPU thread.

```
int main(void) {  
    // call the CUDA kernel from the GPU  
    cuda_hello<<<1,1>>>>();  
  
    // wait for all CUDA kernels to finish  
    cudaDeviceSynchronize();  
    return 0;  
}
```

Hello World

```
#include <cuda.h>

__global__ void cuda_hello(void) {
    // print a character buffer from the GPU!
    printf("Hello, world!")
}

int main(void) {
    // call the CUDA kernel from
    cuda_hello<<<1,1>>>();

    // wait for all CUDA kernels to finish
    cudaDeviceSynchronize();
    return 0;
}
```

Synchronize GPU Threads

Wait for GPU threads to finish. CPU can execute code in the meantime.

Compiling

- Use the Nvidia CUDA C compiler

```
$ nvcc -arch=sm_20 hello.cu
```

```
$ ./a.out
```

```
“Hello, world!”
```

- Different “compute capabilities” on different GPUS.
(Compile targets different hardware levels.)


```
“-arch=sm_53” = Maxwell features
```


Demo

hello.cu

GPU Parallelism

- Shared Memory Environment (e.g. OpenMP)
- Thread organization:



- Set thread / block count using triple chevrons

```
my_kernel<<<nblocks, nthreads_per_block>>>(...)
```

Threads and Block

- Obtain current thread:

```
int tid = threadIdx.x;
```

- Obtain current block:

```
int bid = blockIdx.x;
```

Hello World - Revisited

```
#include <cuda.h>
```

```
__global__ void cuda_hello(void) {  
    int tid = threadIdx.x;  
    printf("Hello, from thread %d\n", tid);  
}
```

Run with 1 block, 4 threads.

```
int main(void) {  
    cuda_hello<<<1,4>>>();  
    cudaDeviceSynchronize();  
    return 0;  
}
```


Demo

hello.cu

Run with 2 blocks and 4 threads per block. “*parallelism coordinate*”

Vector Addition

- Recall three-step memory flow:
 - 1. Copy data from host to device**
 - 2. Compute on device**
 - 3. Copy results back to host**
- Host memory separated from device memory
 - host pointers point to CPU memory
 - device pointers point to GPU memory

vec_add - Device Side

- For now, assume vectors of length N and we spawned 1 block, N threads

```
__global__ void  
vec_add(int* out, int* v, int* w)  
{  
    size_t index = threadIdx.x;  
    out[index] = v[index] + w[index];  
}
```

vec_add - Device Side

- For now, assume vectors of length N and we spawned 1 block, N threads

```
__global__ void  
vec_add(int* out, int* v, int* w)  
{  
    size_t index = threadIdx.x;  
    out[index] = v[index] + w[index];  
}
```

Pointers to device-side memory.

Determine vector index from thread / block indices.

vec_add - Allocating Device Memory

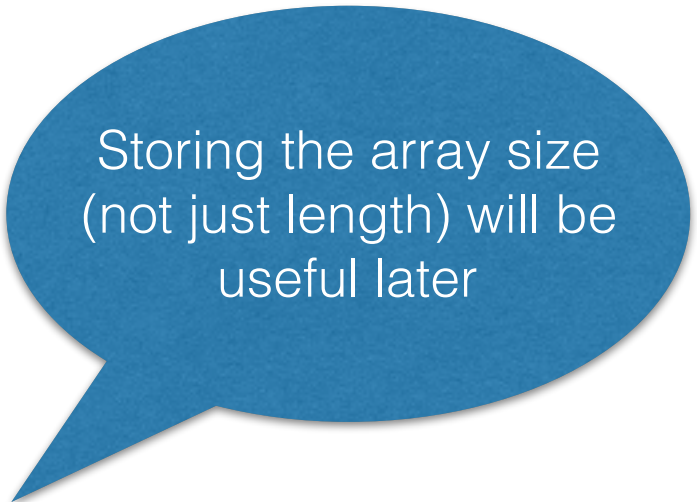
- Device-side analogues of common C memory funs

```
cudaMalloc(&dev_ptr, size);  
cudaMemcpy(dev_ptr, host_ptr, size,  
           direction);  
cudaFree(dev_ptr);
```

Copies data to/from host from/to device. (Depending on dir.)

vec_add - Allocating Device Memory

1) Allocate some host-side data:



Storing the array size
(not just length) will be
useful later

```
int N = 16;  
size_t size = N*sizeof(int);  
int* host_v = (int*) malloc(size);  
int* host_w = (int*) malloc(size);  
  
// populate host_v and host_w with data  
// ...
```

vec_add - Allocating Device Memory

2) Allocate some corresponding device-side data:

```
int* dev_v;  
int* dev_w;  
int* dev_sum;
```

```
// note: cudaMalloc wants ptr to ptr  
cudaMalloc((void**) &dev_v, size);  
cudaMalloc((void**) &dev_w, size);  
cudaMalloc((void**) &dev_out, size);
```

vec_add - Allocating Device Memory

2) Allocate some corresponding device-side data:

```
int* dev_v;  
int* dev_w;  
int* dev_sum;
```

These don't actually point to anywhere in RAM or GPU memory. Just for reference.

```
// note: cudaMalloc wants ptr to ptr  
cudaMalloc((void**) &dev_v, size);  
cudaMalloc((void**) &dev_w, size);  
cudaMalloc((void**) &dev_out, size);
```

vec_add — Allocating Device Memory

- 3) Copy host-side data to device using dev ptrs.

```
cudaMemcpy(dev_v, host_v, size,  
           cudaMemcpyHostToDevice);  
cudaMemcpy(dev_w, host_w, size,  
           cudaMemcpyHostToDevice);
```

- **Remember:** data communication to GPU is explicit (like MPI)

vec_add — Execute GPU Kernel on Device Data

- 4) Launch Kernel on N threads and wait to finish.

```
vec_add<<<1,N>>>(dev_sum,dev_v,dev_w);
```

```
cudaDeviceSynchronize();
```

- **Note:** number of threads spawned = problem size
 - (no need to pass vector length)

vec_add — Copy Results Back to Host

- Remember to free device-side data:

```
cudaMemcpy(host_sum, dev_sum, size,  
           cudaMemcpyDeviceToHost);
```

```
cudaFree(dev_v);  
cudaFree(dev_w);  
cudaFree(dev_sum);
```


```
// ... use results in host_sum ...
```

Demo

vec_add.cu

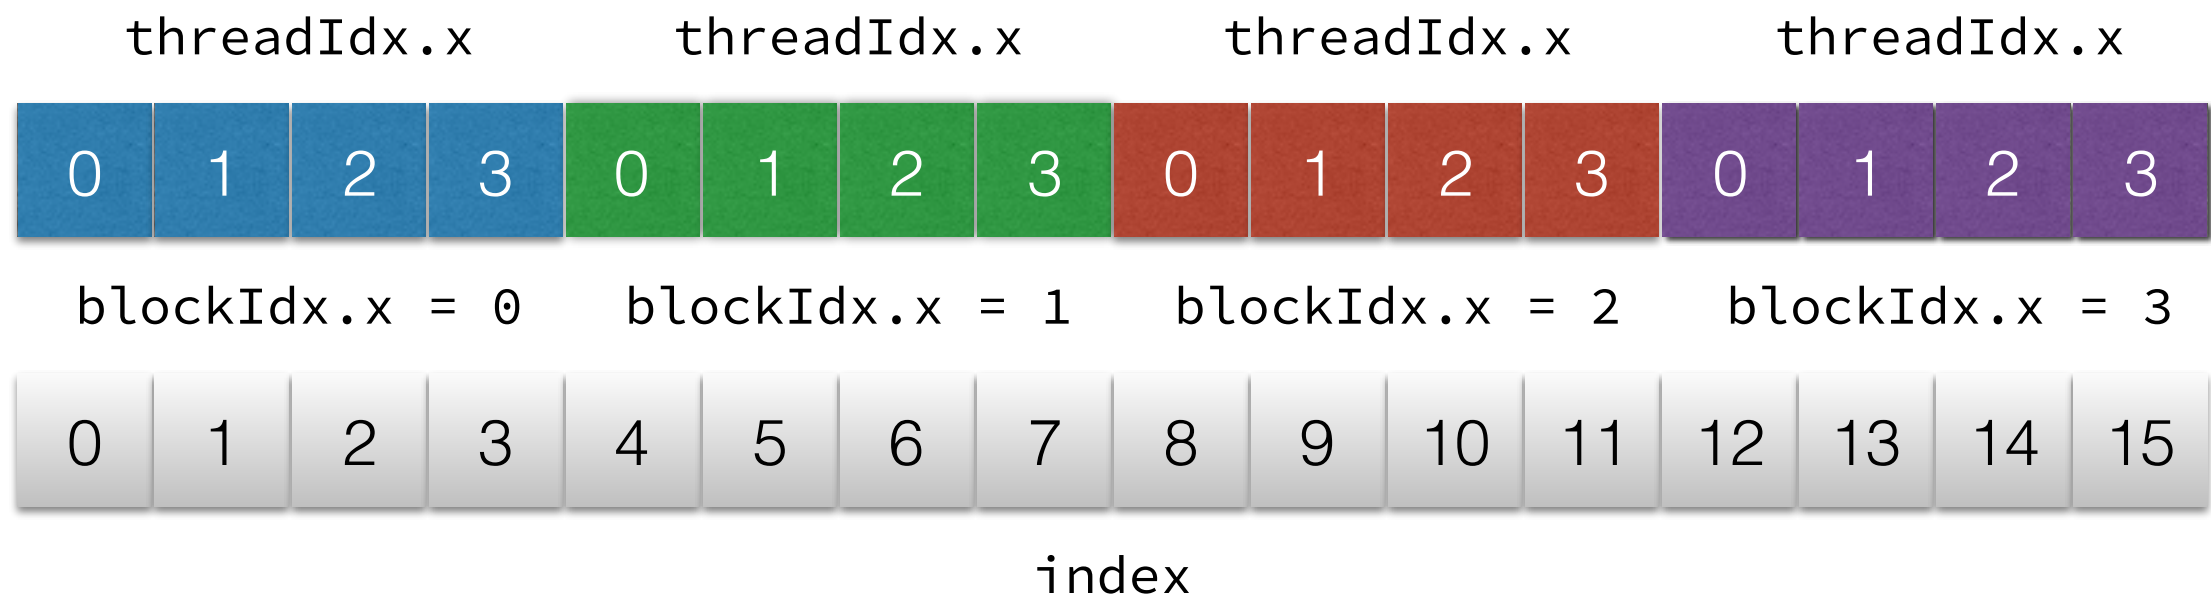
Combining Blocks and Threads

- Each block works on a “chunk” of the vector
- Each block thread adds across an index



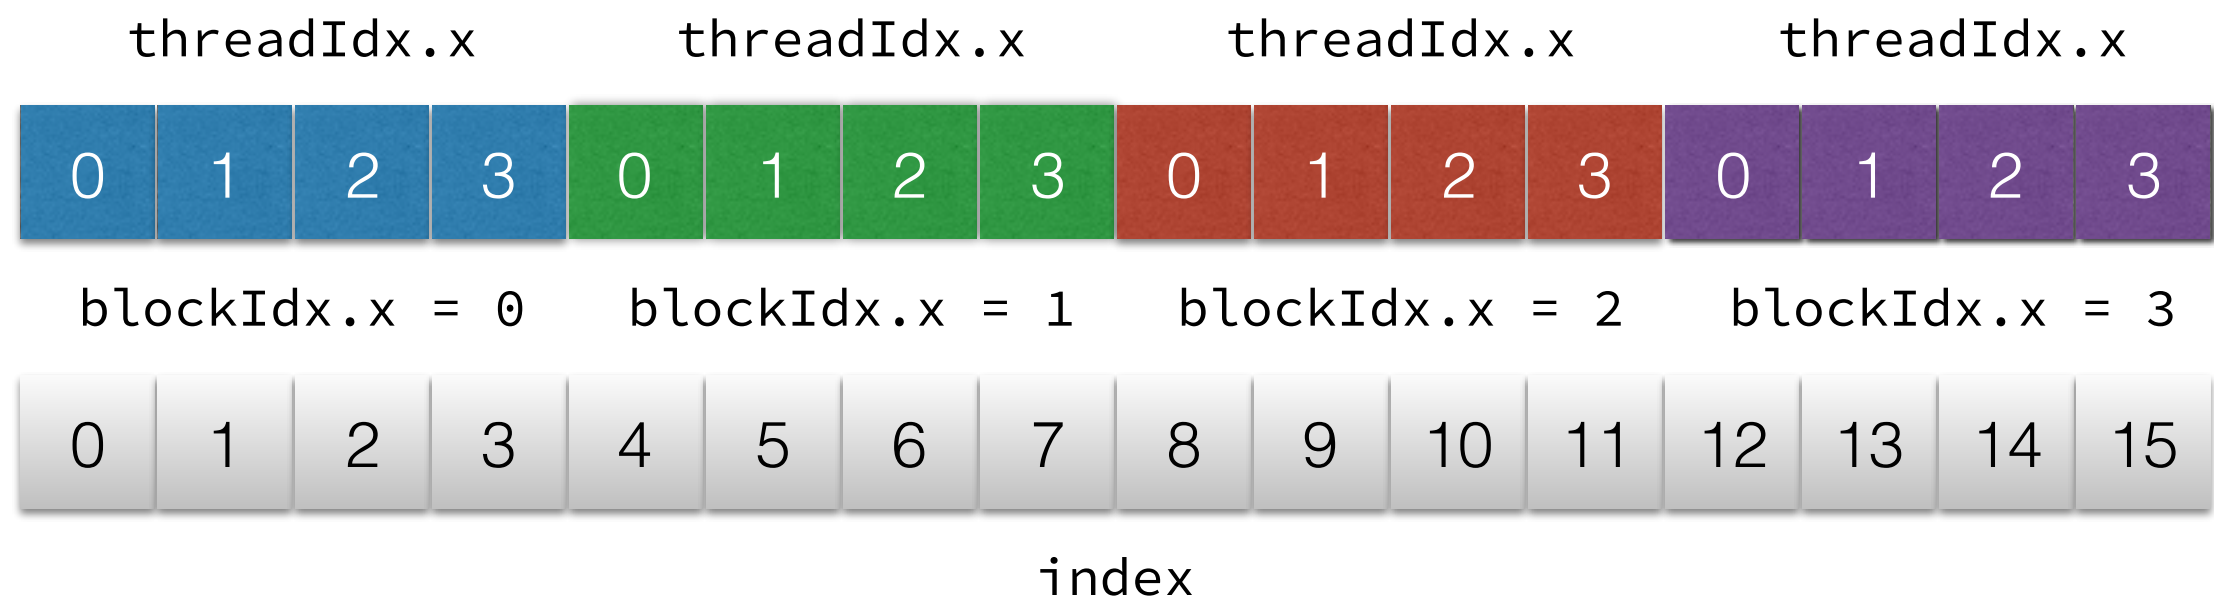
Combining Blocks and Threads

- Compute “global index” = coordinate of vector as a function of `threadIdx.x` and `blockIdx.x`
- `blockDim.x` = number of allocated blocks



Combining Blocks and Threads

```
size_t index;  
index = threadIdx.x +  
        blockIdx.x*blockDim.x;
```



vec_add — Flexible Kernel

- Works for any combination of threads / blocks:

```
__global__ void  
vec_add(int* out, int* v, int* w)  
{  
    size_t index = threadIdx.x +  
                    blockIdx.x * blockDim.x;  
    out[index] = v[index] + w[index];  
}
```

Demo

`vec_add.cu` with variable block/kernel

Why Bother?

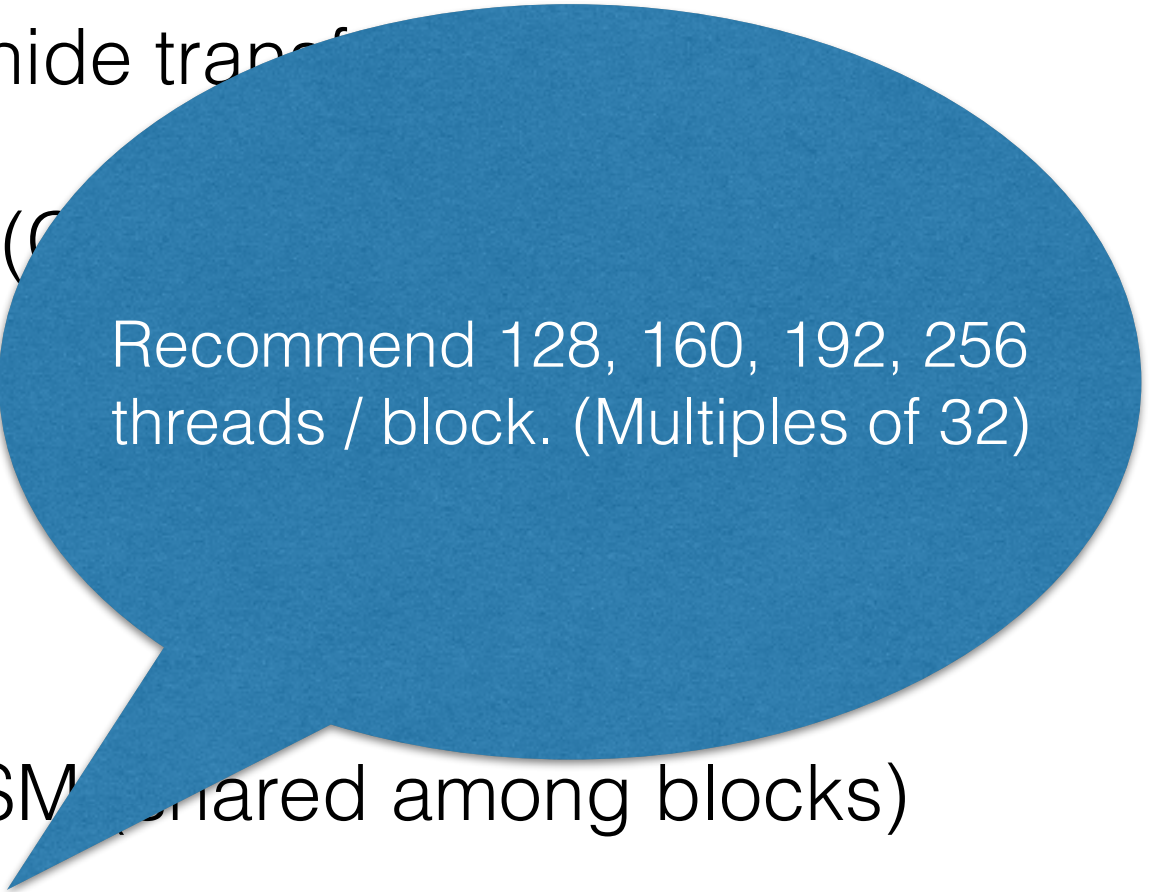
- Within each block threads can
 - **communicate** —
 - *fast “shared memory” between threads within a block (lives in L1 memory)*
 - *otherwise, slow global memory access*
 - **synchronize** — *race cond. and branches*

Why Bother?

- **Latency hiding** — while one block makes mem. request another block can run on SM
 - “block swapping” — hide transfer times
- **Hardware Restrictions** (C2075)
 - 1536 threads / SM
 - 8 blocks / SM
 - total 48 KB shmem / SM (shared among blocks)
 - 1024 threads / block

Why Bother?

- **Latency hiding** — while one block makes mem. request another block can run on SM
 - “block swapping” — hide transfer
- **Hardware Restrictions** (0.5 SM per block)
 - 1536 threads / SM
 - 8 blocks / SM
 - total 48 KB shmem / SM (shared among blocks)
 - 1024 threads / block



Recommend 128, 160, 192, 256 threads / block. (Multiples of 32)

Warps

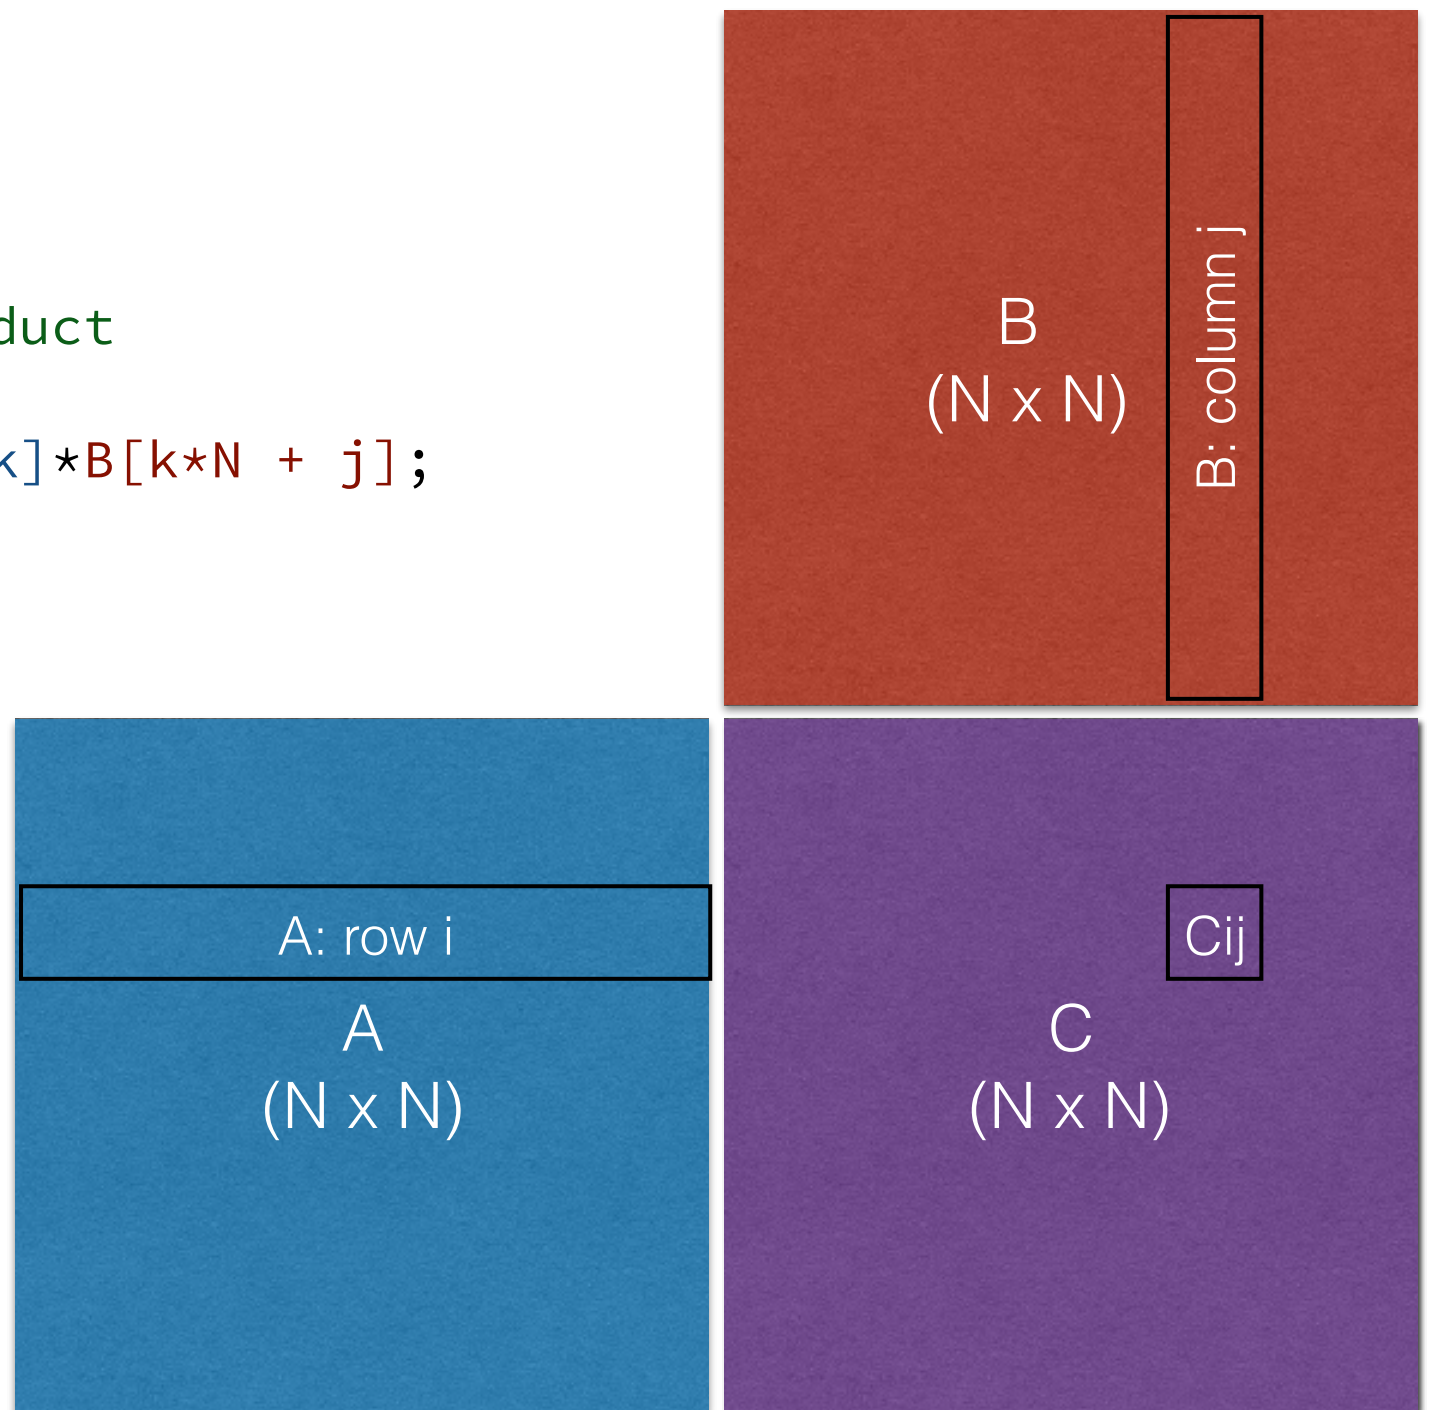
- **Key GPGPU Difference** — each SM splits blocks into “warps” of 32 threads. All threads in a warp execute concurrently.
- warp waits until 32 cores are available on SM
- *if a block contains 48 threads: 32 in one warp 16 in another (+16 no-op threads)*
- performance issues: contiguity, sequential access, aligned access

mat_mul

```
// for each row of C
for (int i=0; i<N; ++i) {
    // for each column of C
    for (int j=0; j<N; ++j) {
        // compute the inner product
        for (int k=0; k<N; ++k)
            C[i*N + j] += A[i*N + k]*B[k*N + j];
    }
}
```

Each thread is assigned
a C_{ij} to compute.

Take advantage of 2D
thread / block arrangement
notation!



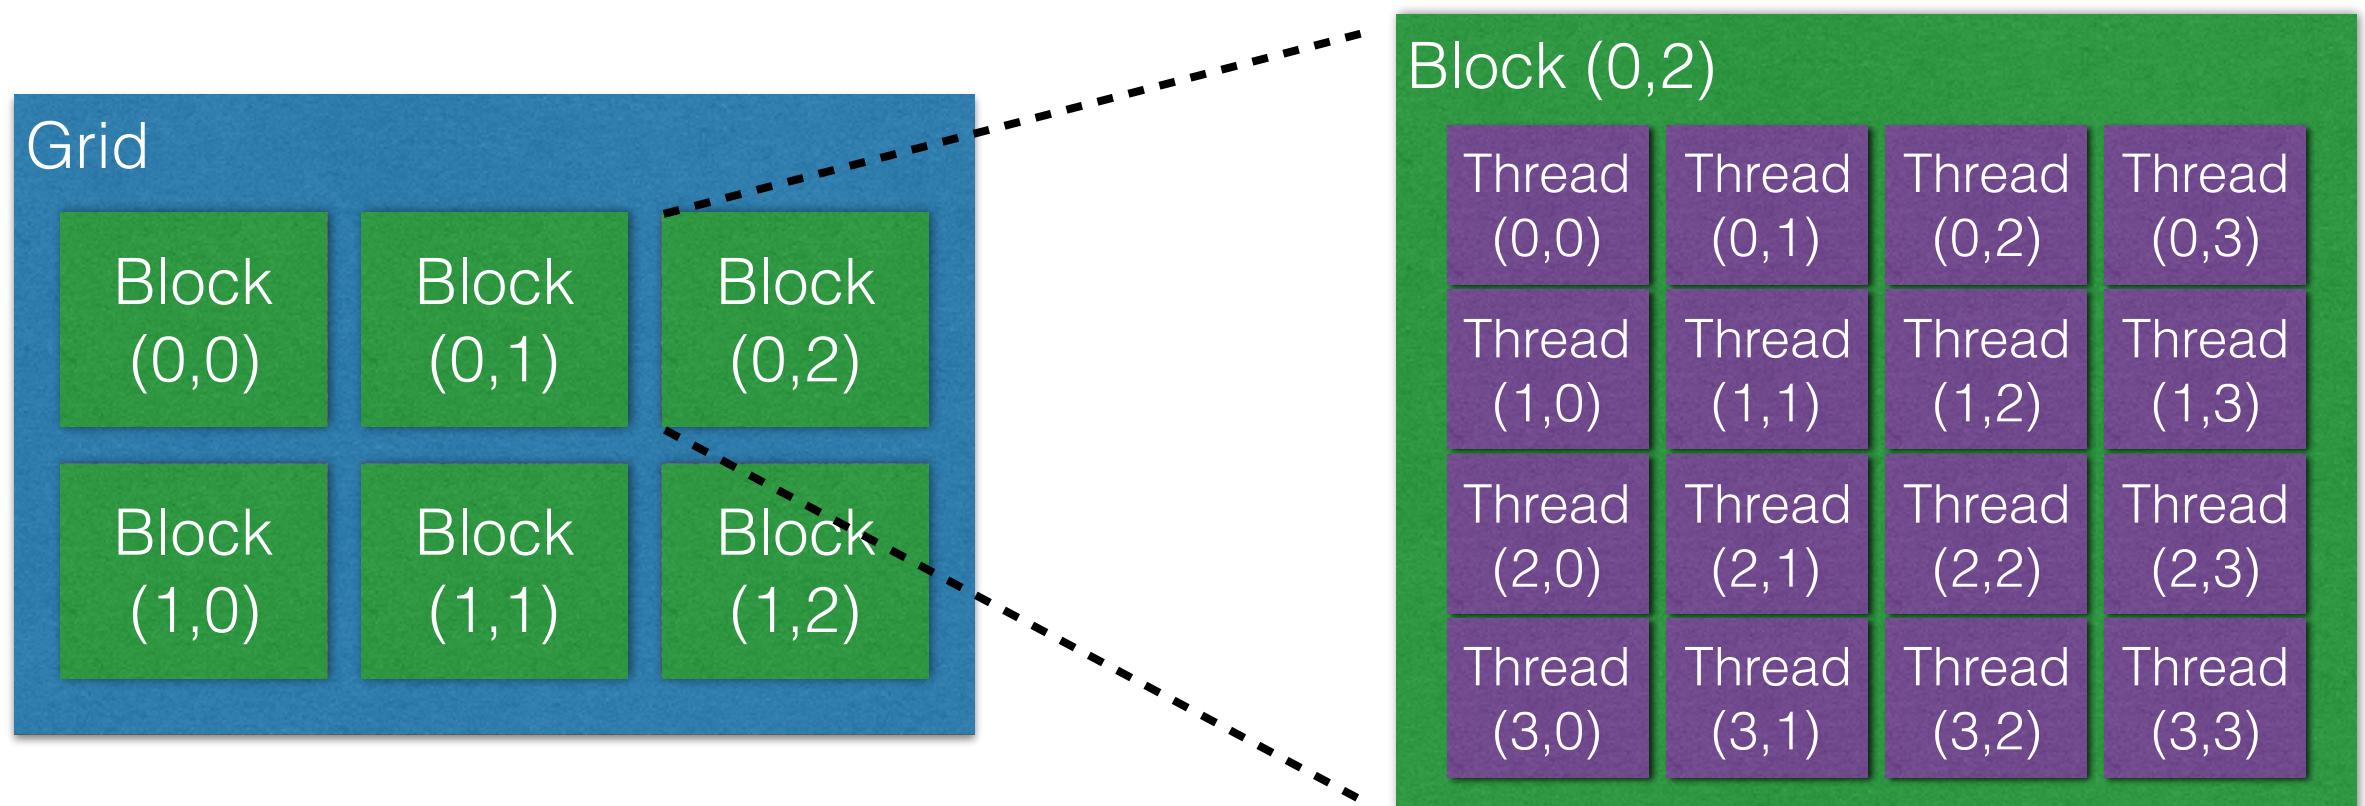
Parallel mat_mul

- Each thread computes C_{ij} — 2D organization of blocks and threads

```
size_t i = threadIdx.y + blockIdx.y*blockDim.y;
```

```
size_t j = threadIdx.x + blockIdx.x*blockDim.x;
```

- Thread at “global index” (i, j) computes $C[i*N+j]$



Creating 2D Grids

- Normal Kernel Call

```
mykernel<<<nblocks, thperblk>>>(...args...);
```

- 2D Grids — if each block contained `thpblk_x * thpblk_y` threads total

```
// specify block size
dim3 dim_block(thpblk_x, thpblk_y);
// specify number of blocks
dim3 dim_grid(nblocks_x, nblocks_y);
mykernel<<<dim_grid, dim_block>>>(...args...);
```

mat_mul — CUDA Kernel

```
__global__ void
mat_mul(double* C, double* A, double* B)
{
    // row#, col#, and row/col length
    size_t i = threadIdx.y + blockIdx.y*blockDim.y;
    size_t j = threadIdx.x + blockIdx.x*blockDim.x;
    size_t N = blockDim.x * gridDim.x;

    // inner product
    for (size_t k=0; k<N; ++k)
        C[i*N+j] = A[i*N+k] * B[k*N+j];
}
```